

AFFR_Log_Reg

May 18, 2020

1 Amazon Fine Food Reviews Analysis

Data Source: <https://www.kaggle.com/snap/amazon-fine-food-reviews>

EDA: <https://nycdatascience.com/blog/student-works/amazon-fine-foods-visualization/>

The Amazon Fine Food Reviews dataset consists of reviews of fine foods from Amazon.

Number of reviews: 568,454 Number of users: 256,059 Number of products: 74,258 Timespan: Oct 1999 - Oct 2012 Number of Attributes/Columns in data: 10

Attribute Information:

1. Id
2. ProductId - unique identifier for the product
3. UserId - unique identifier for the user
4. ProfileName
5. HelpfulnessNumerator - number of users who found the review helpful
6. HelpfulnessDenominator - number of users who indicated whether they found the review helpful or not
7. Score - rating between 1 and 5
8. Time - timestamp for the review
9. Summary - brief summary of the review
10. Text - text of the review

Objective: Given a review, determine whether the review is positive (rating of 4 or 5) or negative (rating of 1 or 2).

[Q] How to determine if a review is positive or negative? [Ans] We could use Score/Rating. A rating of 4 or 5 can be considered as a positive review. A rating of 1 or 2 can be considered as negative one. A review of rating 3 is considered neutral and such reviews are ignored from our analysis. This is an approximate and proxy way of determining the polarity (positivity/negativity) of a review.

2 [1]. Reading Data

2.1 [1.1] Loading the data

The dataset is available in two forms 1. .csv file 2. SQLite Database

In order to load the data, We have used the SQLITE dataset as it is easier to query the data and visualise the data efficiently.

Here as we only want to get the global sentiment of the recommendations (positive or negative), we will purposefully ignore all Scores equal to 3. If the score is above 3, then the recommendation will be set to "positive". Otherwise, it will be set to "negative".

```
[0]: %matplotlib inline
import warnings
warnings.filterwarnings("ignore")

import sqlite3
import pandas as pd
import numpy as np
import nltk
import string
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.feature_extraction.text import TfidfTransformer
from sklearn.feature_extraction.text import TfidfVectorizer
from numpy import random
from sklearn.model_selection import train_test_split

from sklearn.preprocessing import StandardScaler

from sklearn.feature_extraction.text import CountVectorizer
from sklearn.metrics import confusion_matrix
from sklearn import metrics
from sklearn.metrics import roc_curve, auc
from nltk.stem.porter import PorterStemmer

from bs4 import BeautifulSoup

import re
# Tutorial about Python regular expressions: https://pymotw.com/2/re/
import string
from nltk.corpus import stopwords
from nltk.stem import PorterStemmer
from nltk.stem.wordnet import WordNetLemmatizer

from gensim.models import Word2Vec
from gensim.models import KeyedVectors
import pickle

from tqdm import tqdm
import os

from sklearn.metrics import roc_curve, accuracy_score
from sklearn.metrics import precision_score, recall_score
from sklearn.metrics import f1_score, confusion_matrix
```

```
[3]: from google.colab import drive
drive.mount('/content/drive')
```

Go to this URL in a browser: https://accounts.google.com/o/oauth2/auth?client_id=947318989803-6bn6qk8qdgf4n4g3pfee6491hc0brc4i.apps.googleusercontent.com&redirect_uri=urn%3aietf%3awg%3aoauth%3a2.0%3aoob&response_type=code&scope=email%20https%3a%2f%2fwww.googleapis.com%2fauth%2fdocs.test%20https%3a%2f%2fwww.googleapis.com%2fauth%2fdrive%20https%3a%2f%2fwww.googleapis.com%2fauth%2fdrive.photos.readonly%20https%3a%2f%2fwww.googleapis.com%2fauth%2fpeopleapi.readonly

Enter your authorization code:

ûûûûûûûûûûû

Mounted at /content/drive

```
[4]: # using SQLite Table to read data.

con = sqlite3.connect('drive/My Drive/FFRDB/database.sqlite')

# filtering only positive and negative reviews i.e.
# not taking into consideration those reviews with Score=3
# SELECT * FROM Reviews WHERE Score != 3 LIMIT 500000, will give top 500000
→data points
# you can change the number to any other number based on your computing power

# filtered_data = pd.read_sql_query(""" SELECT * FROM Reviews WHERE Score != 3
→LIMIT 500000""", con)
# for tsne assignment you can take 5k data points

filtered_data = pd.read_sql_query(""" SELECT * FROM Reviews WHERE Score != 3
→LIMIT 100000""", con)

# Give reviews with Score>3 a positive rating(1), and reviews with a score<3 a
→negative rating(0).
def partition(x):
    if x < 3:
        return 0
    return 1

#changing reviews with score less than 3 to be positive and vice-versa
actualScore = filtered_data['Score']
positiveNegative = actualScore.map(partition)
filtered_data['Score'] = positiveNegative
print("Number of data points in our data", filtered_data.shape)
filtered_data.head(3)
```

Number of data points in our data (100000, 10)

```
[4]:
```

	Id	...	Text
0	1	...	I have bought several of the Vitality canned d...
1	2	...	Product arrived labeled as Jumbo Salted Peanut...
2	3	...	This is a confection that has been around a fe...

[3 rows x 10 columns]

```
[0]: display = pd.read_sql_query("""
SELECT UserId, ProductId, ProfileName, Time, Score, Text, COUNT(*)
FROM Reviews
GROUP BY UserId
HAVING COUNT(*)>1
""", con)
```

```
[6]: print(display.shape)
display.head()
```

(80668, 7)

```
[6]:
```

	UserId	...	COUNT(*)
0	#oc-R115TNMSPFT9I7	...	2
1	#oc-R11D9D7SHXIJB9	...	3
2	#oc-R11DNU2NBKQ23Z	...	2
3	#oc-R1105J5ZVQE25C	...	3
4	#oc-R12KPBODL2B5ZD	...	2

[5 rows x 7 columns]

```
[7]: display[display['UserId']=='AZY10LLTJ71NX']
```

```
[7]:
```

	UserId	...	COUNT(*)
80638	AZY10LLTJ71NX	...	5

[1 rows x 7 columns]

```
[8]: display['COUNT(*)'].sum()
```

```
[8]: 393063
```

3 [2] Exploratory Data Analysis

3.1 [2.1] Data Cleaning: Deduplication

It is observed (as shown in the table below) that the reviews data had many duplicate entries. Hence it was necessary to remove duplicates in order to get unbiased results for the analysis of the data. Following is an example:

```
[9]: display= pd.read_sql_query("""
SELECT *
FROM Reviews
WHERE Score != 3 AND UserId="AR5J8UI46CURR"
```

```
ORDER BY ProductID
""", con)
display.head()
```

```
[9]:      Id  ...                               Text
0   78445  ...  DELICIOUS WAFERS. I FIND THAT EUROPEAN WAFERS ...
1  138317  ...  DELICIOUS WAFERS. I FIND THAT EUROPEAN WAFERS ...
2  138277  ...  DELICIOUS WAFERS. I FIND THAT EUROPEAN WAFERS ...
3   73791  ...  DELICIOUS WAFERS. I FIND THAT EUROPEAN WAFERS ...
4  155049  ...  DELICIOUS WAFERS. I FIND THAT EUROPEAN WAFERS ...
```

```
[5 rows x 10 columns]
```

As it can be seen above that same user has multiple reviews with same values for HelpfulnessNumerator, HelpfulnessDenominator, Score, Time, Summary and Text and on doing analysis it was found that ProductId=B000HDOPZG was Loacker Quadratini Vanilla Wafer Cookies, 8.82-Ounce Packages (Pack of 8) ProductId=B000HDL1RQ was Loacker Quadratini Lemon Wafer Cookies, 8.82-Ounce Packages (Pack of 8) and so on

It was inferred after analysis that reviews with same parameters other than ProductId belonged to the same product just having different flavour or quantity. Hence in order to reduce redundancy it was decided to eliminate the rows having same parameters.

The method used for the same was that we first sort the data according to ProductId and then just keep the first similar product review and delete the others. for eg. in the above just the review for ProductId=B000HDL1RQ remains. This method ensures that there is only one representative for each product and deduplication without sorting would lead to possibility of different representatives still existing for the same product.

```
[0]: #Sorting data according to ProductId in ascending order
sorted_data=filtered_data.sort_values('ProductId', axis=0, ascending=True,
→inplace=False, kind='quicksort', na_position='last')
```

```
[11]: #Deduplication of entries
final=sorted_data.
→drop_duplicates(subset={"UserId","ProfileName","Time","Text"}, keep='first',
→inplace=False)
final.shape
```

```
[11]: (87775, 10)
```

```
[12]: final.sort_values('Time',inplace=True)
print(final.head(5))
```

```
      Id  ...                               Text
70688  76882  ...  I bought a few of these after my apartment was...
1146   1245  ...  This was a really good idea and the final prod...
1145   1244  ...  I just received my shipment and could hardly w...
28086  30629  ...  Nothing against the product, but it does bothe...
28087  30630  ...  I love this stuff. It is sugar-free so it does...
```

```
[5 rows x 10 columns]
```

```
[13]: #Checking to see how much % of data still remains
(final['Id'].size*1.0)/(filtered_data['Id'].size*1.0)*100
```

[13]: 87.775

Observation:- It was also seen that in two rows given below the value of HelpfulnessNumerator is greater than HelpfulnessDenominator which is not practically possible hence these two rows too are removed from calculations

```
[14]: display= pd.read_sql_query("""
SELECT *
FROM Reviews
WHERE Score != 3 AND Id=44737 OR Id=64422
ORDER BY ProductID
""", con)

display.head()
```

```
[14]:      Id  ...                               Text
0  64422  ...  My son loves spaghetti so I didn't hesitate or...
1  44737  ...  It was almost a 'love at first bite' - the per...
```

[2 rows x 10 columns]

```
[0]: final=final[final.HelpfulnessNumerator<=final.HelpfulnessDenominator]
```

```
[16]: #Before starting the next phase of preprocessing lets see the number of entries
      ↪left
print(final.shape)

#How many positive and negative reviews are present in our dataset?
final['Score'].value_counts()
```

(87773, 10)

```
[16]: 1    73592
      0    14181
      Name: Score, dtype: int64
```

4 [3] Preprocessing

4.1 [3.1]. Preprocessing Review Text

Now that we have finished deduplication our data requires some preprocessing before we go on further with analysis and making the prediction model.

Hence in the Preprocessing phase we do the following in the order below:-

1. Begin by removing the html tags
2. Remove any punctuations or limited set of special characters like , or . or # etc.
3. Check if the word is made up of english letters and is not alpha-numeric

4. Check to see if the length of the word is greater than 2 (as it was researched that there is no adjective in 2-letters)
5. Convert the word to lowercase
6. Remove Stopwords
7. Finally Snowball Stemming the word (it was observed to be better than Porter Stemming)

After which we collect the words used to describe positive and negative reviews

```
[17]: # printing some random reviews
sent_0 = final['Text'].values[0]
print(sent_0)
print("="*50)

sent_1000 = final['Text'].values[1000]
print(sent_1000)
print("="*50)

sent_1500 = final['Text'].values[1500]
print(sent_1500)
print("="*50)

sent_4900 = final['Text'].values[4900]
print(sent_4900)
print("="*50)
```

I bought a few of these after my apartment was infested with fruit flies. After only a few hours, the trap had "attracted" many flies and within a few days they were practically gone. This may not be a long term solution, but if flies are driving you crazy, consider buying this. One caution- the surface is very sticky, so try to avoid touching it.

=====

we use this as the base, then besides the chicken, we will also add pasta, spices, veggies, or whatever we have around to make quick cheesy meals

=====

My dogs just love this food. The service is always fast and reliable.

=====

I am amazed by how well this tea works to relieve my chronic congestion and recurring sinus problems. And it's not just a "quick" fix either -- its therapeutic effects last for hours. I was a bit worried the tea would be a bit too "licorice-y" since one of its main ingredients is licorice root, but the fragrance and taste are mild and incredibly soothing. If you think this package of six boxes is too much, you'll be happily proven wrong ... I would stock my entire garage with this tea if I could!

=====

```
[18]: # remove urls from text python: https://stackoverflow.com/a/40823105/4084039
sent_0 = re.sub(r"http\S+", "", sent_0)
sent_1000 = re.sub(r"http\S+", "", sent_1000)
```

```
sent_150 = re.sub(r"http\S+", "", sent_1500)
sent_4900 = re.sub(r"http\S+", "", sent_4900)

print(sent_0)
```

I bought a few of these after my apartment was infested with fruit flies. After only a few hours, the trap had "attracted" many flies and within a few days they were practically gone. This may not be a long term solution, but if flies are driving you crazy, consider buying this. One caution- the surface is very sticky, so try to avoid touching it.

```
[19]: # https://stackoverflow.com/questions/16206380/
      ↪python-beautifulsoup-how-to-remove-all-tags-from-an-element
from bs4 import BeautifulSoup

soup = BeautifulSoup(sent_0, 'lxml')
text = soup.get_text()
print(text)
print("="*50)

soup = BeautifulSoup(sent_1000, 'lxml')
text = soup.get_text()
print(text)
print("="*50)

soup = BeautifulSoup(sent_1500, 'lxml')
text = soup.get_text()
print(text)
print("="*50)

soup = BeautifulSoup(sent_4900, 'lxml')
text = soup.get_text()
print(text)
```

I bought a few of these after my apartment was infested with fruit flies. After only a few hours, the trap had "attracted" many flies and within a few days they were practically gone. This may not be a long term solution, but if flies are driving you crazy, consider buying this. One caution- the surface is very sticky, so try to avoid touching it.

=====

we use this as the base, then besides the chicken, we will also add pasta, spices, veggies, or whatever we have around to make quick cheesy meals

=====

My dogs just love this food. The service is always fast and reliable.

=====

I am amazed by how well this tea works to relieve my chronic congestion and recurring sinus problems. And it's not just a "quick" fix either -- its

therapeutic effects last for hours. I was a bit worried the tea would be a bit too "licorice-y" since one of its main ingredients is licorice root, but the fragrance and taste are mild and incredibly soothing. If you think this package of six boxes is too much, you'll be happily proven wrong ... I would stock my entire garage with this tea if I could!

```
[0]: # https://stackoverflow.com/a/47091490/4084039
import re

def decontracted(phrase):
    # specific
    phrase = re.sub(r"won't", "will not", phrase)
    phrase = re.sub(r"can't", "can not", phrase)

    # general
    phrase = re.sub(r"n't", " not", phrase)
    phrase = re.sub(r"\ 're", " are", phrase)
    phrase = re.sub(r"\ 's", " is", phrase)
    phrase = re.sub(r"\ 'd", " would", phrase)
    phrase = re.sub(r"\ 'll", " will", phrase)
    phrase = re.sub(r"\ 't", " not", phrase)
    phrase = re.sub(r"\ 've", " have", phrase)
    phrase = re.sub(r"\ 'm", " am", phrase)
    return phrase
```

```
[21]: sent_1500 = decontracted(sent_1500)
print(sent_1500)
print("="*50)
```

My dogs just love this food. The service is always fast and reliable.
=====

```
[22]: #remove words with numbers python: https://stackoverflow.com/a/18082370/4084039
sent_0 = re.sub("\S*\d\S*", "", sent_0).strip()
print(sent_0)
```

I bought a few of these after my apartment was infested with fruit flies. After only a few hours, the trap had "attracted" many flies and within a few days they were practically gone. This may not be a long term solution, but if flies are driving you crazy, consider buying this. One caution- the surface is very sticky, so try to avoid touching it.

```
[23]: #remove spacial character: https://stackoverflow.com/a/5843547/4084039
sent_1500 = re.sub('[^A-Za-z0-9]+', ' ', sent_1500)
print(sent_1500)
```

My dogs just love this food The service is always fast and reliable

```
[0]: # https://gist.github.com/sebleier/554280
# we are removing the words from the stop words list: 'no', 'nor', 'not'
# <br /><br /> ==> after the above steps, we are getting "br br"
# we are including them into stop words list
# instead of <br /> if we have <br/> these tags would have been removed in the 1st
→ step

stopwords= set(['br', 'the', 'i', 'me', 'my', 'myself', 'we', 'our', 'ours',
→ 'ourselves', 'you', "you're", "you've",\
    "you'll", "you'd", 'your', 'yours', 'yourself', 'yourselves', 'he',
→ 'him', 'his', 'himself', \
    'she', "she's", 'her', 'hers', 'herself', 'it', "it's", 'its',
→ 'itself', 'they', 'them', 'their',\
    'theirs', 'themselves', 'what', 'which', 'who', 'whom', 'this',
→ 'that', "that'll", 'these', 'those', \
    'am', 'is', 'are', 'was', 'were', 'be', 'been', 'being', 'have',
→ 'has', 'had', 'having', 'do', 'does', \
    'did', 'doing', 'a', 'an', 'the', 'and', 'but', 'if', 'or',
→ 'because', 'as', 'until', 'while', 'of', \
    'at', 'by', 'for', 'with', 'about', 'against', 'between', 'into',
→ 'through', 'during', 'before', 'after',\
    'above', 'below', 'to', 'from', 'up', 'down', 'in', 'out', 'on',
→ 'off', 'over', 'under', 'again', 'further',\
    'then', 'once', 'here', 'there', 'when', 'where', 'why', 'how',
→ 'all', 'any', 'both', 'each', 'few', 'more',\
    'most', 'other', 'some', 'such', 'only', 'own', 'same', 'so',
→ 'than', 'too', 'very', \
    's', 't', 'can', 'will', 'just', 'don', "don't", 'should',
→ "should've", 'now', 'd', 'll', 'm', 'o', 're', \
    've', 'y', 'ain', 'aren', "aren't", 'couldn', "couldn't", 'didn',
→ "didn't", 'doesn', "doesn't", 'hadn',\
    "hadn't", 'hasn', "hasn't", 'haven', "haven't", 'isn', "isn't",
→ 'ma', 'mightn', "mightn't", 'mustn',\
    "mustn't", 'needn', "needn't", 'shan', "shan't", 'shouldn',
→ "shouldn't", 'wasn', "wasn't", 'weren', "weren't", \
    'won', "won't", 'wouldn', "wouldn't"])
```

```
[25]: # Combining all the above students
from tqdm import tqdm
preprocessed_reviews = []
# tqdm is for printing the status bar
for sentence in tqdm(final['Text'].values):
    sentence = re.sub(r"http\S+", "", sentence)
    sentence = BeautifulSoup(sentence, 'lxml').get_text()
    sentence = decontracted(sentence)
    sentence = re.sub(r"\S*d\S*", "", sentence).strip()
```

```

sentence = re.sub('[^A-Za-z]+', ' ', sentence)
# https://gist.github.com/sebleier/554280
sentence = ' '.join(e.lower() for e in sentence.split() if e.lower() not in
→stopwords)
preprocessed_reviews.append(sentence.strip())

```

100%|| 87773/87773 [00:30<00:00, 2914.19it/s]

[26]: preprocessed_reviews[1500]

[26]: 'dogs love food service always fast reliable'

[3.2] Preprocessing Review Summary

[0]: *## Similarly you can do preprocessing for review summary also.*

5 [4] Featurization

5.1 [4.4.1] loading tfidf and w2v pickles

```

[0]: import pickle
import os

dbfile1 = open('/content/drive/My Drive/FFRDB/tfidf.pkl', 'rb')
tfidf_sent_vectors = pickle.load(dbfile1)

dbfile2 = open('/content/drive/My Drive/FFRDB/sent_vectors.pkl', 'rb')
sent_vectors= pickle.load(dbfile2)

```

6 [5] Assignment 5: Apply Logistic Regression

Apply Logistic Regression on these feature sets

SET 1:Review text, preprocessed one converted into vectors

SET 2:Review text, preprocessed one converted into vectors

SET 3:Review text, preprocessed one converted into vectors

SET 4:Review text, preprocessed one converted into vectors

Hyper paramter tuning (find best hyper parameters corresponding the algorithm that

Find the best hyper parameter which will give the maximum <a href='https://www.applieaicom

Find the best hyper paramter using k-fold cross validation or simple cross validation data

Use gridsearch cv or randomsearch cv or you can also write your own for loops to do this t


```

</li>
<br>
<li><strong>Pertubation Test</strong>
  <ul>
<li>Get the weights W after fit your model with the data X i.e Train data.</li>
<li>Add a noise to the X ( $X' = X + e$ ) and get the new data set X' (if X is a sparse
    matrix,  $X.data += e$ )

<li>Fit the model again on data X' and get the weights W'</li>
<li>Add a small eps value(to eliminate the divisible by zero error) to W and W i.e

     $W = W + 10^{-6}$  and  $W' = W' + 10^{-6}$ 

<li>Now find the % change between W and W' ( $| (W - W') / (W) | * 100$ )</li>
<li>Calculate the 0th, 10th, 20th, 30th, ...100th percentiles, and observe any sudden rise in
<li> Ex: consider your 99th percentile is 1.3 and your 100th percentiles are 34.6, there is su
    <li> Print the feature names whose % change is more than a threshold x(in our example
  </ul>
</li>
<br>
<li><strong>Sparsity</strong>
  <ul>
<li>Calculate sparsity on weight vector obtained after using L1 regularization</li>
  </ul>
</li>
<br>
<font color='red'>NOTE: Do sparsity and multicollinearity for any one of the vectorizers.
<br>
<br>
<li><strong>Feature importance</strong>
  <ul>
<li>Get top 10 important features for both positive and negative classes separately.</li>
  </ul>
</li>
<br>
<li><strong>Feature engineering</strong>
  <ul>
<li>To increase the performance of your model, you can also experiment with with feature engineering
    <ul>
<li>Taking length of reviews as another feature.</li>
<li>Considering some features from review summary as well.</li>
    </ul>
  </ul>
</li>
<br>
<li><strong>Representation of results</strong>
  <ul>
<li>You need to plot the performance of model both on train data and cross validation data for
<img src='train_cv_auc.JPG' width=300px></li>

```

```

<li>Once after you found the best hyper parameter, you need to train your model with it, and f
<img src='train_test_auc.JPG' width=300px></li>
<li>Along with plotting ROC curve, you need to print the <a href='https://www.appliedaicourse.
<img src='confusion_matrix.png' width=300px></li>
</ul>
</li>
<br>
<li><strong>Conclusion</strong>
<ul>
<li>You need to summarize the results at the end of the notebook, summarize it in the table for
<img src='summary.JPG' width=400px>
</li>
</ul>

```

Note: Data Leakage

1. There will be an issue of data-leakage if you vectorize the entire data and then split it into train/cv/test.
2. To avoid the issue of data-leakag, make sure to split your data first and then vectorize it.
3. While vectorizing your data, apply the method fit_transform() on you train data, and apply the method transform() on cv/test data.
4. For more details please go through this link.

7 Applying Logistic Regression

7.1 [5.1] Logistic Regression on BOW, SET 1

7.1.1 [5.1.1] Applying Logistic Regression with L1 regularization on BOW, SET 1

```

[99]: #generating random alpha values between 10^-5 to 10^5
from numpy import random
c=[]
#removing 0 as value from alpha
def remove_zero(a):
    while True:
        if 0 in a:
            a.pop(0)
            continue
        break
#removing dupliacte
def remove_dup(a):
    for i,j in enumerate(a):
        if i!=(len(a)-1):
            if j == a[i+1]:
                a.pop(i)

#generating alpha values one at a time
def generator(a):

```

```

x,y=random.randint(-3,3),random.randint(0,9)
z=round(y*10**x,abs(x))
a.extend([z])
return a.sort()

while len(c)!=40:
    generator(c)
    remove_dup(c)
    remove_zero(c)

# c=[1000,100,10,1,0.1,0.01,0.001,0.0001]
print(c)
print(len(c))

```

[0.001, 0.002, 0.003, 0.004, 0.005, 0.006, 0.007, 0.008, 0.01, 0.02, 0.04, 0.05, 0.06, 0.1, 0.4, 0.5, 0.6, 0.8, 1, 2, 3, 5, 6, 7, 8, 10, 20, 30, 50, 60, 70, 80, 100, 200, 300, 400, 500, 600, 700, 800]

40

```

[102]: from sklearn.feature_extraction.text import CountVectorizer
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler

X_train, X_test, y_train, y_test = _
    →train_test_split(preprocessed_reviews,final['Score'].values,test_size=0.
    →3,random_state=0)

vectorizer = CountVectorizer(ngram_range=(1,2), min_df=10,max_features=5000)
vectorizer.fit(X_train)
X_train = vectorizer.transform(X_train)
X_test = vectorizer.transform(X_test)

ss = StandardScaler(with_mean = False)
X_train = ss.fit_transform(X_train)
X_test = ss.transform(X_test)

print(X_train.shape)
print(X_test.shape)

```

(61441, 5000)
(26332, 5000)

```
[103]: param = {'C':c}
print(param)

from sklearn.model_selection import GridSearchCV
from sklearn.linear_model import LogisticRegression
l1= LogisticRegression(penalty='l1', solver='liblinear')
temp_gscv=
    ↪GridSearchCV(l1,param,cv=5,verbose=5,n_jobs=-1,scoring='roc_auc',return_train_score=True)
temp_gscv.fit(X_train,y_train)
temp_gs = pd.DataFrame.from_dict(temp_gscv.cv_results_)
```


```
{'C': [0.001, 0.002, 0.003, 0.004, 0.005, 0.006, 0.007, 0.008, 0.01, 0.02, 0.04,
0.05, 0.06, 0.1, 0.4, 0.5, 0.6, 0.8, 1, 2, 3, 5, 6, 7, 8, 10, 20, 30, 50, 60,
70, 80, 100, 200, 300, 400, 500, 600, 700, 800]}
```

Fitting 5 folds for each of 40 candidates, totalling 200 fits

```
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 2 concurrent workers.
[Parallel(n_jobs=-1)]: Done 14 tasks      | elapsed:    4.1s
[Parallel(n_jobs=-1)]: Done 68 tasks      | elapsed:   26.5s
[Parallel(n_jobs=-1)]: Done 158 tasks     | elapsed:   4.7min
[Parallel(n_jobs=-1)]: Done 200 out of 200 | elapsed:   7.1min finished
```

```
[104]: train_auc= temp_gs['mean_train_score']
cv_auc= temp_gs['mean_test_score']
plt.plot(np.log(c),train_auc,label='Train AUC')
plt.plot(np.log(c),cv_auc ,label='CV AUC')

# plt.scatter(param['n_neighbors'],train_auc,label='Train AUC')
# plt.scatter(param['n_neighbors'], cv_auc, label='cv AUC')
plt.grid()
plt.title('AUC vs C plot')
plt.xlabel("log(C): hyperparameter")
plt.ylabel("AUC")
plt.legend()
plt.show()
plt.clf()
plt.cla()
plt.close()
```



[105]: *#finding the best CV score first then using the one which is least distant then
→its AUC counter part to avoid using Dumb model.*

```
from scipy.signal import argrelextrema
import numpy as np
x = np.array(train_auc)
y = np.array(cv_auc)
local_max=()
#finding local maximas of CV
local_max_i=argrelextrema(y, np.greater)
l=list(i for i in np.nditer(local_max_i))

diff=x-y
# diff between CV and Test AUC at the local maxima
local_diff=list(diff[i] for i in l)
local_c=list(c[i] for i in l)
print(f'all local differences {local_diff}')
print(f'all local max C {local_c}')

for i in np.nditer(np.argmin(local_diff)):
    v=i
    break
print(f'best cv score to use = {y[l[v]]}')
```



```
best_c=c[l[v]]
print(f'best C to use = {c[l[v]]}')
```

```
all local differences [0.02530679255826418]
all local max C [0.02]
best cv score to use = 0.9488122858028978
best C to use = 0.02
```


```
[0]: from sklearn.linear_model import LogisticRegression

from sklearn.metrics import
    →accuracy_score, confusion_matrix, f1_score, precision_score, recall_score, roc_curve,
    →auc

lr = LogisticRegression(penalty='l1', C=best_c, solver='liblinear')
lr.fit(X_train, y_train)
y_pred_tr = lr.predict_proba(X_train)
y_pred_ts = lr.predict_proba(X_test)
y_pred_ts=y_pred_ts[:,1]
y_pred_tr = y_pred_tr[:,1]
```

```
[107]: train_fpr, train_tpr, tr_thresholds = roc_curve(y_train, y_pred_tr)
test_fpr, test_tpr, te_thresholds = roc_curve(y_test, y_pred_ts)

plt.plot(train_fpr, train_tpr, label="train AUC =" + str(auc(train_fpr,
    →train_tpr)))
plt.plot(test_fpr, test_tpr, label="test AUC =" + str(auc(test_fpr, test_tpr)))
plt.xlabel("fpr")
plt.ylabel("tpr")
plt.title('ROC curve for ' + str(best_c) + 'as C')
plt.legend()
plt.grid()
plt.show()
```



```
[108]: # This section of code where ever implemented is taken from sample kNN python_
        ↳notebook

def find_best_threshold(threshold, fpr, tpr):
    t = threshold[np.argmax(tpr*(1-fpr))]
    # (tpr*(1-fpr)) will be maximum if your fpr is very low and tpr is very_
    ↳high
    print("the maximum value of tpr*(1-fpr)", max(tpr*(1-fpr)), "for_
    ↳threshold", np.round(t,3))
    return t

def predict_with_best_t(proba, threshold):
    predictions = []
    for i in proba:
        if i>=threshold:
            predictions.append(1)
        else:
            predictions.append(0)
    return predictions

print('test')
best_ts_thres = find_best_threshold(te_thresholds, test_fpr, test_tpr)
```

```
print('train')
best_tr_thres = find_best_threshold(tr_thresholds, train_fpr, train_tpr)
```

test

the maximum value of $tpr \cdot (1 - fpr)$ 0.7896444699590777 for threshold 0.821


train

the maximum value of $tpr \cdot (1 - fpr)$ 0.8482000151187687 for threshold 0.805

```
[109]: print('train Confusion Matrix')
cm2 = pd.DataFrame(confusion_matrix(y_train, predict_with_best_t(y_pred_tr,
    ↳best_tr_thres)), range(2),range(2))
sns.heatmap(cm2, annot=True,fmt='g')
```

train Confusion Matrix


```
[109]: <matplotlib.axes._subplots.AxesSubplot at 0x7f3a4a3fb3c8>
```



```
[110]: print('Test Confusion Matrix')
cm2 = pd.DataFrame(confusion_matrix(y_test, predict_with_best_t(y_pred_ts,
    ↳best_ts_thres)), range(2),range(2))
sns.heatmap(cm2, annot=True,fmt='g')
```

Test Confusion Matrix

[110]: <matplotlib.axes._subplots.AxesSubplot at 0x7f3a4a80d588>



```
[111]: acc=accuracy_score(y_test, predict_with_best_t(y_pred_ts, best_ts_thres))*100
ps=precision_score(y_test, predict_with_best_t(y_pred_ts, best_ts_thres))*100
rc=recall_score(y_test, predict_with_best_t(y_pred_ts, best_ts_thres))*100
f1=f1_score(y_test, predict_with_best_t(y_pred_ts, best_ts_thres))*100

print("Accuracy on test set: %0.2f%%"%(acc))
print("Precision on test set: %0.2f%%"%(ps))
print("recall score on test set: %0.2f%%"%(rc))
print("f1 score on test set: %0.2f%%"%(f1))
```

Accuracy on test set: 89.36%
Precision on test set: 97.52%
recall score on test set: 89.59%
f1 score on test set: 93.39%

[5.1.1.1] Calculating sparsity on weight vector obtained using L1 regularization on BOW, SET 1

```
[112]: non_zero_features=np.count_nonzero(lr.coef_)
print(f"Number of features with non-zero weights :{non_zero_features}")
print(f"Number of features with zero weights :{5000-non_zero_features}")
```

Number of features with non-zero weights :2361
Number of features with zero weights :2639

7.1.2 [5.1.2] Applying Logistic Regression with L2 regularization on BOW, SET 1

[113]: *#generating random alpha values between 10^{-5} to 10^5*

```
from numpy import random
c=[]
#removing 0 as value from alpha
def remove_zero(a):
    while True:
        if 0 in a:
            a.pop(0)
            continue
        break
#removing duplicate
def remove_dup(a):
    for i,j in enumerate(a):
        if i!=(len(a)-1):
            if j == a[i+1]:
                a.pop(i)

#generating alpha values one at a time
def generator(a):
    x,y=random.randint(-3,3),random.randint(0,9)
    z=round(y*10**x,abs(x))
    a.extend([z])
    return a.sort()

while len(c)!=40:
    generator(c)
    remove_dup(c)
    remove_zero(c)

# c=[1000,100,10,1,0.1,0.01,0.001,0.0001]
print(c)
print(len(c))
```

```
[0.001, 0.002, 0.003, 0.004, 0.006, 0.01, 0.02, 0.03, 0.04, 0.05, 0.06, 0.07,
0.08, 0.1, 0.2, 0.3, 0.4, 0.5, 0.6, 0.7, 1, 2, 3, 4, 5, 6, 7, 8, 10, 20, 30, 50,
60, 80, 200, 400, 500, 600, 700, 800]
40
```

[114]: `from sklearn.feature_extraction.text import CountVectorizer`
`from sklearn.model_selection import train_test_split`
`from sklearn.preprocessing import StandardScaler`

```

X_train, X_test, y_train, y_test = ␣
    →train_test_split(preprocessed_reviews,final['Score'].values,test_size=0.
    →3,random_state=0)

vectorizer = CountVectorizer(ngram_range=(1,2), min_df=10, max_features=5000)
vectorizer.fit(X_train)
X_train = vectorizer.transform(X_train)
X_test = vectorizer.transform(X_test)

ss = StandardScaler(with_mean = False)
X_train = ss.fit_transform(X_train)
X_test = ss.transform(X_test)

print(X_train.shape)
print(X_test.shape)

```

```

(61441, 5000)
(26332, 5000)

```

```

[115]: param = {'C':c}
print(param)

from sklearn.model_selection import GridSearchCV
from sklearn.linear_model import LogisticRegression
l1= LogisticRegression(penalty='l2')
temp_gscv=␣
    →GridSearchCV(l1,param,cv=5,verbose=5,n_jobs=-1,scoring='roc_auc',return_train_score=True)
temp_gscv.fit(X_train,y_train)
temp_gs = pd.DataFrame.from_dict(temp_gscv.cv_results_)

```

```

{'C': [0.001, 0.002, 0.003, 0.004, 0.006, 0.01, 0.02, 0.03, 0.04, 0.05, 0.06,
0.07, 0.08, 0.1, 0.2, 0.3, 0.4, 0.5, 0.6, 0.7, 1, 2, 3, 4, 5, 6, 7, 8, 10, 20,
30, 50, 60, 80, 200, 400, 500, 600, 700, 800]}

```

Fitting 5 folds for each of 40 candidates, totalling 200 fits

```

[Parallel(n_jobs=-1)]: Using backend LokyBackend with 2 concurrent workers.
[Parallel(n_jobs=-1)]: Done 14 tasks      | elapsed:    6.5s
[Parallel(n_jobs=-1)]: Done 68 tasks      | elapsed:   50.8s
[Parallel(n_jobs=-1)]: Done 158 tasks     | elapsed:   2.1min
[Parallel(n_jobs=-1)]: Done 200 out of 200 | elapsed:   2.7min finished

```

```

[116]: train_auc= temp_gs['mean_train_score']
cv_auc= temp_gs['mean_test_score']
plt.plot(np.log(c),train_auc,label='Train AUC')
plt.plot(np.log(c),cv_auc ,label='CV AUC')


# plt.scatter(param['n_neighbors'],train_auc,label='Train AUC')

```

```

# plt.scatter(param['n_neighbors'], cv_auc, label='cv AUC')
plt.grid()
plt.title('AUC vs C plot')
plt.xlabel("log(C): hyperparameter")
plt.ylabel("AUC")
plt.legend()
plt.show()
plt.clf()
plt.cla()
plt.close()

```



[117]: *#finding the best CV score first then using the one which is least distant then
→its AUC counter part to avoid using Dumb model.*

```

from scipy.signal import argrelextrema
import numpy as np
x = np.array(train_auc)
y = np.array(cv_auc)
local_max=()
#finding local maximas of CV
local_max_i=argrelextrema(y, np.greater)
l=list(i for i in np.nditer(local_max_i))

```

```

diff=x-y
# diff between CV and Test AUC at the local maxima
local_diff=list(diff[i] for i in l)
local_c=list(c[i] for i in l)
print(f'all local differences {local_diff}')
print(f'all local max C {local_c}')

for i in np.nditer(np.argmin(local_diff)):
    v=i
    break
print(f'best cv score to use = {y[l[v]]}')
best_c=c[l[v]]
print(f'best C to use = {c[l[v]]}')

```

all local differences [0.07553428709873644, 0.07560867286356143,
0.07563993861404084, 0.07561146158188525, 0.07570301744801444]
all local max C [7, 10, 60, 200, 600]
best cv score to use = 0.9157965703904208
best C to use = 7

```

[0]: from sklearn.linear_model import LogisticRegression

from sklearn.metrics import
    →accuracy_score, confusion_matrix, f1_score, precision_score, recall_score, roc_curve,
    →auc

lr = LogisticRegression(penalty='l2', C=best_c)
lr.fit(X_train, y_train)
y_pred_tr = lr.predict_proba(X_train)
y_pred_ts = lr.predict_proba(X_test)
y_pred_ts=y_pred_ts[:,1]
y_pred_tr = y_pred_tr[:,1]


```

```

[119]: train_fpr, train_tpr, tr_thresholds = roc_curve(y_train, y_pred_tr)
test_fpr, test_tpr, te_thresholds = roc_curve(y_test, y_pred_ts)

plt.plot(train_fpr, train_tpr, label="train AUC =" + str(auc(train_fpr,
    →train_tpr)))
plt.plot(test_fpr, test_tpr, label="test AUC =" + str(auc(test_fpr, test_tpr)))
plt.xlabel("fpr")
plt.ylabel("tpr")
plt.title('ROC curve for ' + str(best_c) + ' as C')
plt.legend()
plt.grid()
plt.show()

```

[120]: *# This section of code where ever implemented is taken from sample kNN python_*
→notebook

```
def find_best_threshold(threshold, fpr, tpr):
    t = threshold[np.argmax(tpr*(1-fpr))]
    # (tpr*(1-fpr)) will be maximum if your fpr is very low and tpr is very_
    →high
    print("the maximum value of tpr*(1-fpr)", max(tpr*(1-fpr)), "for_
    →threshold", np.round(t,3))
    return t

def predict_with_best_t(proba, threshold):
    predictions = []
    for i in proba:
        if i>=threshold:
            predictions.append(1)
        else:
            predictions.append(0)
    return predictions

print('test')
best_ts_thres = find_best_threshold(te_thresholds, test_fpr, test_tpr)
```

```
print('train')
best_tr_thres = find_best_threshold(tr_thresholds, train_fpr, train_tpr)
```

test

the maximum value of $tpr \cdot (1 - fpr)$ 0.738226035005531 for threshold 0.926


train

the maximum value of $tpr \cdot (1 - fpr)$ 0.8855549937161102 for threshold 0.818

```
[121]: print('train Confusion Matrix')
cm2 = pd.DataFrame(confusion_matrix(y_train, predict_with_best_t(y_pred_tr,
    ↳best_tr_thres)), range(2),range(2))
sns.heatmap(cm2, annot=True,fmt='g')
```

train Confusion Matrix


```
[121]: <matplotlib.axes._subplots.AxesSubplot at 0x7f3a4ab6c908>
```



```
[122]: print('Test Confusion Matrix')
cm2 = pd.DataFrame(confusion_matrix(y_test, predict_with_best_t(y_pred_ts,
    ↳best_ts_thres)), range(2),range(2))
sns.heatmap(cm2, annot=True,fmt='g')
```

Test Confusion Matrix

[122]: <matplotlib.axes._subplots.AxesSubplot at 0x7f3a4b47c358>



```
[123]: acc=accuracy_score(y_test, predict_with_best_t(y_pred_ts, best_ts_thres))*100
ps=precision_score(y_test, predict_with_best_t(y_pred_ts, best_ts_thres))*100
rc=recall_score(y_test, predict_with_best_t(y_pred_ts, best_ts_thres))*100
f1=f1_score(y_test, predict_with_best_t(y_pred_ts, best_ts_thres))*100

print("Accuracy on test set: %0.2f%%"%(acc))
print("Precision on test set: %0.2f%%"%(ps))
print("recall score on test set: %0.2f%%"%(rc))
print("f1 score on test set: %0.2f%%"%(f1))
```

Accuracy on test set: 86.84%
Precision on test set: 96.71%
recall score on test set: 87.28%
f1 score on test set: 91.76%

```
[0]: count=0
value=[]
for i in lr.coef_.reshape(-1,1):
    count+=1
    value.extend(i)

x=vectorizer.get_feature_names()

log_features= pd.DataFrame({'feature_name':x,'value':value})
```

[5.1.2.1] Performing pertubation test (multicollinearity check) on BOW, SET 1

```
[125]: # Please write all the code with proper documentation
```

```
epsilon = 0.00002
X_train.data = X_train.data +(epsilon,)
X_train.shape
```

```
[125]: (61441, 5000)
```

```
[0]: from sklearn.linear_model import LogisticRegression

from sklearn.metrics import accuracy_score, confusion_matrix, f1_score, precision_score, recall_score, roc_curve, auc


xlr = LogisticRegression(penalty='l2', C=best_c)
xlr.fit(X_train, y_train)
y_pred_tr = xlr.predict_proba(X_train)
y_pred_ts = xlr.predict_proba(X_test)
y_pred_ts=y_pred_ts[:,1]
y_pred_tr = y_pred_tr[:,1]
```

```
[0]: w1=lr.coef_+0.000001
w2=xlr.coef_+0.000001
```

```
[0]: p=abs((w1[0]-w2[0])/w1[0]*100)
p=list(i for i in p)
```

```
[129]: change=pd.DataFrame({'change_per':(p), 'feature':x}, index=range(count))
change.sort_values(by='change_per',ascending=False, inplace=True)
# change=change[change['change_per']>0]
plt.plot(range(change.shape[0]),np.log(change['change_per']))
```

```
[129]: [ <matplotlib.lines.Line2D at 0x7f3a4b1036d8>]
```



```
[0]: '''
Algorithm to find elbow of a graph is taken from the following questionnaire on
→Stackoverflow
#####
https://stackoverflow.com/questions/2018178/
→finding-the-best-trade-off-point-on-a-curve
#####
'''
def elbow_finder(x_values):
    import numpy as np
    import numpy.matlib
    nPoints = len(x_values)
    allCoord = np.vstack((range(nPoints), x_values)).T
    np.array([range(nPoints), x_values])
    firstPoint = allCoord[0]
    lineVec = allCoord[-1] - allCoord[0]
    lineVecNorm = lineVec / np.sqrt(np.sum(lineVec**2))
    vecFromFirst = allCoord - firstPoint
    scalarProduct = np.sum(vecFromFirst * np.matlib.repmat(lineVecNorm, nPoints, 1), axis=1)
    vecFromFirstParallel = np.outer(scalarProduct, lineVecNorm)
    vecToLine = vecFromFirst - vecFromFirstParallel
    distToLine = np.sqrt(np.sum(vecToLine ** 2, axis=1))
    return np.argmax(distToLine)
```

```
[131]: x=elbow_finder(change['change_per'])
print(f'all the index before {x} is multicollinear as this is the elbow point')
```

all the index before 36 is multicollinear as this is the elbow point

```
[132]: change.iloc[:x]
```

```
[132]:      change_per      feature
3363  3467.691841      pour
3466  2274.404191  protein bars
4108  2180.684584      spring
4149  1253.551494      steak
1451   389.218805  favorite coffee
2445   216.970657  list ingredients
1313   213.763843      equivalent
1496   147.240872      final
3733   141.610713      royal
3577   127.106397  really love
4230   121.101915      sugar
424    113.218241      boston
620    100.178705      cat not
3540    97.777559      rate
3485    95.590974      purchase
120    93.335267      although
3885    83.012935      shaking
4382    82.639976      tea great
3212    80.022064      paste
3543    77.451867      rating
3807    69.629883  science diet
1385    61.950954      experienced
1731    59.158479      get rid
2265    56.865670      kinda
2211    56.684352      jar
423    55.647775      boring
3119    53.003202  orange juice
1599    51.988146      food
4563    51.584418      treats dog
1633    47.873415  found product
1583    47.695039      flour
3267    45.317203      pg
27     44.956753      across
1835    44.293774      goodies
1393    42.078029      extra
3100    38.150695      ones not
```

7.1.3 [5.1.3] Feature Importance on BOW, SET 1

[5.1.3.1] Top 10 important features of positive class from SET 1

```
[133]: log_features.sort_values(by = ['value'], ascending=False).head(10)
```

```
[133]:
```

	feature_name	value
1872	great	1.022723
3797	say enough	0.901300
1072	delicious	0.818206
3698	right size	0.808981
1791	good	0.749260
3695	right amount	0.705956
325	best	0.654930
2054	highly recommend	0.624402
2834	no bitterness	0.582643
3249	perfect	0.573723

[5.1.3.2] Top 10 important features of negative class from SET 1

```
[134]: log_features.sort_values(by = ['value'], ascending=True).head(10)
```

```
[134]:
```

	feature_name	value
2867	not	-0.969615
1145	disappointed	-0.449436
1240	earth best	-0.404937
4906	worst	-0.380955
1776	goat milk	-0.379761
3432	product	-0.368179
3003	not worth	-0.364820
4411	terrible	-0.353169
772	coco	-0.335525
243	bad	-0.330583

7.2 [5.2] Logistic Regression on TFIDF, SET 2

7.2.1 [5.2.1] Applying Logistic Regression with L1 regularization on TFIDF, SET 2

```
[177]: #generating random alpha values between 10^-5 to 10^5
from numpy import random
c=[]
#removing 0 as value from alpha
def remove_zero(a):
    while True:
        if 0 in a:
            a.pop(0)
            continue
        break
#removing duplicate
def remove_dup(a):
    for i,j in enumerate(a):
        if i!=(len(a)-1):
```

```

        if j == a[i+1]:
            a.pop(i)

#generating alpha values one at a time
def generator(a):
    x,y=random.randint(-3,3),random.randint(0,9)
    z=round(y*10**x,abs(x))
    a.extend([z])
    return a.sort()

while len(c)!=40:
    generator(c)
    remove_dup(c)
    remove_zero(c)

# c=[1000,100,10,1,0.1,0.01,0.001,0.0001]
print(c)
print(len(c))

```

[0.001, 0.003, 0.004, 0.007, 0.02, 0.03, 0.04, 0.05, 0.06, 0.07, 0.08, 0.1, 0.2, 0.3, 0.4, 0.5, 0.6, 0.7, 0.8, 1, 2, 3, 4, 5, 6, 7, 8, 10, 30, 50, 70, 80, 100, 200, 300, 400, 500, 600, 700, 800]
40

```

[175]: from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler

X_train, X_test, y_train, y_test = _
    →train_test_split(preprocessed_reviews,final['Score'].values,test_size=0.
    →3,random_state=0)

vectorizer = TfidfVectorizer(ngram_range=(1,2), min_df=10, max_features=5000)
vectorizer.fit(X_train)
X_train = vectorizer.transform(X_train)
X_test = vectorizer.transform(X_test)

ss = StandardScaler(with_mean = False)
X_train = ss.fit_transform(X_train)
X_test = ss.transform(X_test)

print(X_train.shape)
print(X_test.shape)

```



```
(61441, 5000)
(26332, 5000)
```

```
[176]: param = {'C':c}
print(param)

from sklearn.model_selection import GridSearchCV
from sklearn.linear_model import LogisticRegression
l1= LogisticRegression(penalty='l1', solver='liblinear')
temp_gscv=
    ↳GridSearchCV(l1,param,cv=5,verbose=5,n_jobs=-1,scoring='roc_auc',return_train_score=True)
temp_gscv.fit(X_train,y_train)
temp_gs = pd.DataFrame.from_dict(temp_gscv.cv_results_)
```

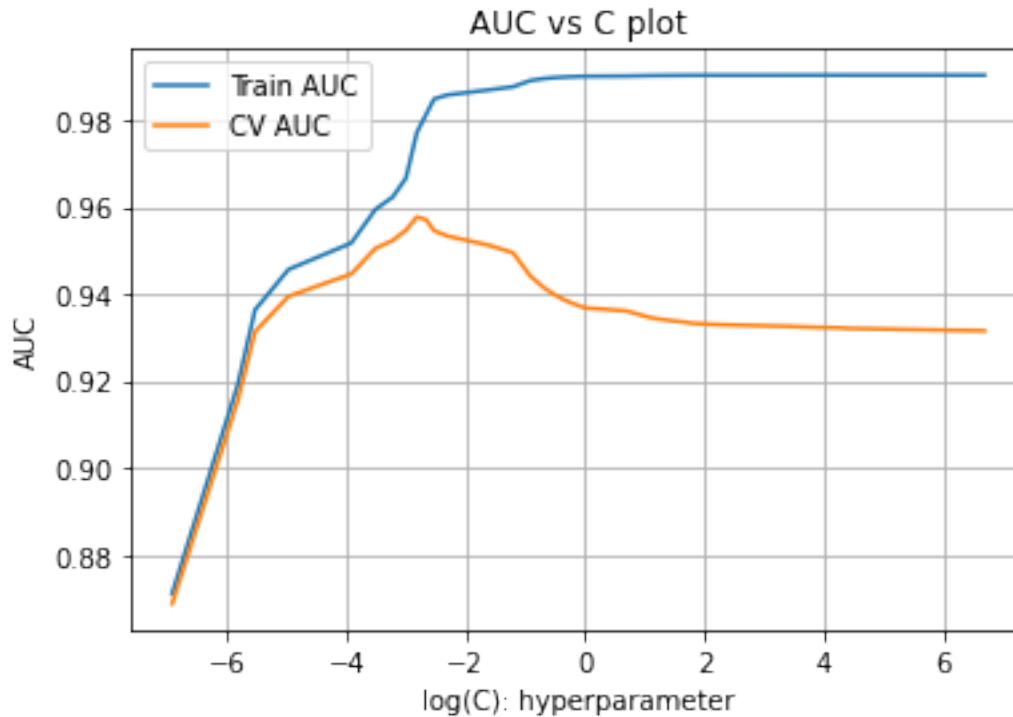
```
{'C': [0.001, 0.002, 0.003, 0.004, 0.005, 0.007, 0.008, 0.01, 0.02, 0.03, 0.05,
0.06, 0.08, 0.1, 0.2, 0.3, 0.4, 0.5, 0.6, 0.8, 1, 2, 3, 4, 6, 7, 8, 10, 20, 40,
50, 80, 100, 200, 300, 400, 500, 600, 700, 800]}
```

Fitting 5 folds for each of 40 candidates, totalling 200 fits

```
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 2 concurrent workers.
[Parallel(n_jobs=-1)]: Done 14 tasks      | elapsed:    3.7s
[Parallel(n_jobs=-1)]: Done 68 tasks      | elapsed:   37.3s
[Parallel(n_jobs=-1)]: Done 158 tasks     | elapsed:   3.5min
[Parallel(n_jobs=-1)]: Done 200 out of 200 | elapsed:   5.1min finished
```

```
[178]: train_auc= temp_gs['mean_train_score']
cv_auc= temp_gs['mean_test_score']
plt.plot(np.log(c),train_auc,label='Train AUC')
plt.plot(np.log(c),cv_auc ,label='CV AUC')

# plt.scatter(param['n_neighbors'],train_auc,label='Train AUC')
# plt.scatter(param['n_neighbors'], cv_auc, label='cv AUC')
plt.grid()
plt.title('AUC vs C plot')
plt.xlabel("log(C): hyperparameter")
plt.ylabel("AUC")
plt.legend()
plt.show()
plt.clf()
plt.cla()
plt.close()
```



[179]: *#finding the best CV score first then using the one which is least distant then_*
→its AUC counter part to avoid using Dumb model.

```
from scipy.signal import argrelextrema
import numpy as np
x = np.array(train_auc)
y = np.array(cv_auc)
local_max=()
#finding local maximas of CV
local_max_i=argrelextrema(y, np.greater)
l=list(i for i in np.nditer(local_max_i))

diff=x-y
# diff between CV and Test AUC at the local maxima
local_diff=list(diff[i] for i in l)
local_c=list(c[i] for i in l)
print(f'all local differences {local_diff}')
print(f'all local max C {local_c}')

for i in np.nditer(np.argmin(local_diff)):
    v=i
    break
print(f'best cv score to use = {y[l[v]]}')
```

```
best_c=c[l[v]]
print(f'best C to use = {c[l[v]]}')
```

```
all local differences [0.019429676063882795]
all local max C [0.06]
best cv score to use = 0.9577134597053177
best C to use = 0.06
```


```
[0]: from sklearn.linear_model import LogisticRegression

from sklearn.metrics import
    →accuracy_score, confusion_matrix, f1_score, precision_score, recall_score, roc_curve,
    →auc

lr = LogisticRegression(penalty='l1', C=best_c, solver='liblinear')
lr.fit(X_train, y_train)
y_pred_tr = lr.predict_proba(X_train)
y_pred_ts = lr.predict_proba(X_test)
y_pred_ts=y_pred_ts[:,1]
y_pred_tr = y_pred_tr[:,1]
```

```
[181]: train_fpr, train_tpr, tr_thresholds = roc_curve(y_train, y_pred_tr)
test_fpr, test_tpr, te_thresholds = roc_curve(y_test, y_pred_ts)

plt.plot(train_fpr, train_tpr, label="train AUC =" + str(auc(train_fpr,
    →train_tpr)))
plt.plot(test_fpr, test_tpr, label="test AUC =" + str(auc(test_fpr, test_tpr)))
plt.xlabel("fpr")
plt.ylabel("tpr")
plt.title('ROC curve for ' + str(best_c) + 'as C')
plt.legend()
plt.grid()
plt.show()
```



[182]: *# This section of code where ever implemented is taken from sample kNN python_*
→notebook

```
def find_best_threshold(threshold, fpr, tpr):
    t = threshold[np.argmax(tpr*(1-fpr))]
    # (tpr*(1-fpr)) will be maximum if your fpr is very low and tpr is very_
    →high
    print("the maximum value of tpr*(1-fpr)", max(tpr*(1-fpr)), "for_
    →threshold", np.round(t,3))
    return t

def predict_with_best_t(proba, threshold):
    predictions = []
    for i in proba:
        if i>=threshold:
            predictions.append(1)
        else:
            predictions.append(0)
    return predictions

print('test')
best_ts_thres = find_best_threshold(te_thresholds, test_fpr, test_tpr)
```

```
print('train')
best_tr_thres = find_best_threshold(tr_thresholds, train_fpr, train_tpr)
```

test

the maximum value of $tpr \cdot (1 - fpr)$ 0.7838396244130014 for threshold 0.865


train

the maximum value of $tpr \cdot (1 - fpr)$ 0.8759119333204912 for threshold 0.806

```
[183]: print('train Confusion Matrix')
cm2 = pd.DataFrame(confusion_matrix(y_train, predict_with_best_t(y_pred_tr,
    ↳best_tr_thres)), range(2),range(2))
sns.heatmap(cm2, annot=True,fmt='g')
```

train Confusion Matrix


```
[183]: <matplotlib.axes._subplots.AxesSubplot at 0x7f3a4df4c438>
```



```
[184]: print('Test Confusion Matrix')
cm2 = pd.DataFrame(confusion_matrix(y_test, predict_with_best_t(y_pred_ts,
    ↳best_ts_thres)), range(2),range(2))
sns.heatmap(cm2, annot=True,fmt='g')
```

Test Confusion Matrix

[184]: <matplotlib.axes._subplots.AxesSubplot at 0x7f3a4aed0668>



```
[185]: acc=accuracy_score(y_test, predict_with_best_t(y_pred_ts, best_ts_thres))*100
ps=precision_score(y_test, predict_with_best_t(y_pred_ts, best_ts_thres))*100
rc=recall_score(y_test, predict_with_best_t(y_pred_ts, best_ts_thres))*100
f1=f1_score(y_test, predict_with_best_t(y_pred_ts, best_ts_thres))*100

print("Accuracy on test set: %0.2f%%"%(acc))
print("Precision on test set: %0.2f%%"%(ps))
print("recall score on test set: %0.2f%%"%(rc))
print("f1 score on test set: %0.2f%%"%(f1))
```

```
Accuracy on test set: 88.47%
Precision on test set: 97.59%
recall score on test set: 88.44%
f1 score on test set: 92.79%
```

```
[186]: non_zero_features=np.count_nonzero(lr.coef_)
print(f"Number of features with non-zero weights :{non_zero_features}")
print(f"Number of features with zero weights :{5000-non_zero_features}")
```

```
Number of features with non-zero weights :3854
Number of features with zero weights :1146
```

7.2.2 [5.2.2] Applying Logistic Regression with L2 regularization on TFIDE, SET 2

[147]: *#generating random alpha values between 10^{-5} to 10^5*

```
from numpy import random
c=[]
#removing 0 as value from alpha
def remove_zero(a):
    while True:
        if 0 in a:
            a.pop(0)
            continue
        break
#removing dupliacte
def remove_dup(a):
    for i,j in enumerate(a):
        if i!=(len(a)-1):
            if j == a[i+1]:
                a.pop(i)

#generating alpha values one at a time
def generator(a):
    x,y=random.randint(-3,3),random.randint(0,9)
    z=round(y*10**x,abs(x))
    a.extend([z])
    return a.sort()

while len(c)!=40:
    generator(c)
    remove_dup(c)
    remove_zero(c)

# c=[1000,100,10,1,0.1,0.01,0.001,0.0001]
print(c)
print(len(c))
```

```
[0.002, 0.004, 0.005, 0.006, 0.007, 0.008, 0.01, 0.02, 0.03, 0.04, 0.06, 0.07,
0.08, 0.1, 0.2, 0.3, 0.4, 0.5, 0.6, 1, 2, 3, 4, 5, 6, 7, 10, 30, 40, 50, 60, 70,
80, 100, 200, 300, 400, 500, 600, 800]
40
```

```
[67]: from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
```

```

X_train, X_test, y_train, y_test = u
    →train_test_split(preprocessed_reviews,final['Score'].values,test_size=0.
    →3,random_state=0)

vectorizer = TfidfVectorizer(ngram_range=(1,2), min_df=10, max_features=5000)
vectorizer.fit(X_train)
X_train = vectorizer.transform(X_train)
X_test = vectorizer.transform(X_test)

ss = StandardScaler(with_mean = False)
X_train = ss.fit_transform(X_train)
X_test = ss.transform(X_test)

print(X_train.shape)
print(X_test.shape)

```

```

(61441, 5000)
(26332, 5000)

```

```

[68]: param = {'C':c}
print(param)

from sklearn.model_selection import GridSearchCV
from sklearn.linear_model import LogisticRegression
l1= LogisticRegression(penalty='l2')
temp_gscv=u
    →GridSearchCV(l1,param,cv=5,verbose=5,n_jobs=-1,scoring='roc_auc',return_train_score=True)
temp_gscv.fit(X_train,y_train)
temp_gs = pd.DataFrame.from_dict(temp_gscv.cv_results_)

```

```

{'C': [0.001, 0.003, 0.004, 0.005, 0.006, 0.007, 0.01, 0.02, 0.03, 0.05, 0.06,
0.07, 0.08, 0.1, 0.2, 0.3, 0.5, 0.6, 0.7, 0.8, 1, 2, 3, 4, 5, 6, 7, 8, 10, 20,
30, 50, 70, 100, 200, 300, 400, 500, 600, 800]}

```

Fitting 5 folds for each of 40 candidates, totalling 200 fits

```

[Parallel(n_jobs=-1)]: Using backend LokyBackend with 2 concurrent workers.
[Parallel(n_jobs=-1)]: Done 14 tasks      | elapsed:    7.8s
[Parallel(n_jobs=-1)]: Done 68 tasks      | elapsed:   48.9s
[Parallel(n_jobs=-1)]: Done 158 tasks     | elapsed:   2.1min
[Parallel(n_jobs=-1)]: Done 200 out of 200 | elapsed:   2.7min finished

```

```

[69]: train_auc= temp_gs['mean_train_score']
cv_auc= temp_gs['mean_test_score']
plt.plot(np.log(c),train_auc,label='Train AUC')
plt.plot(np.log(c),cv_auc ,label='CV AUC')

# plt.scatter(param['n_neighbors'],train_auc,label='Train AUC')

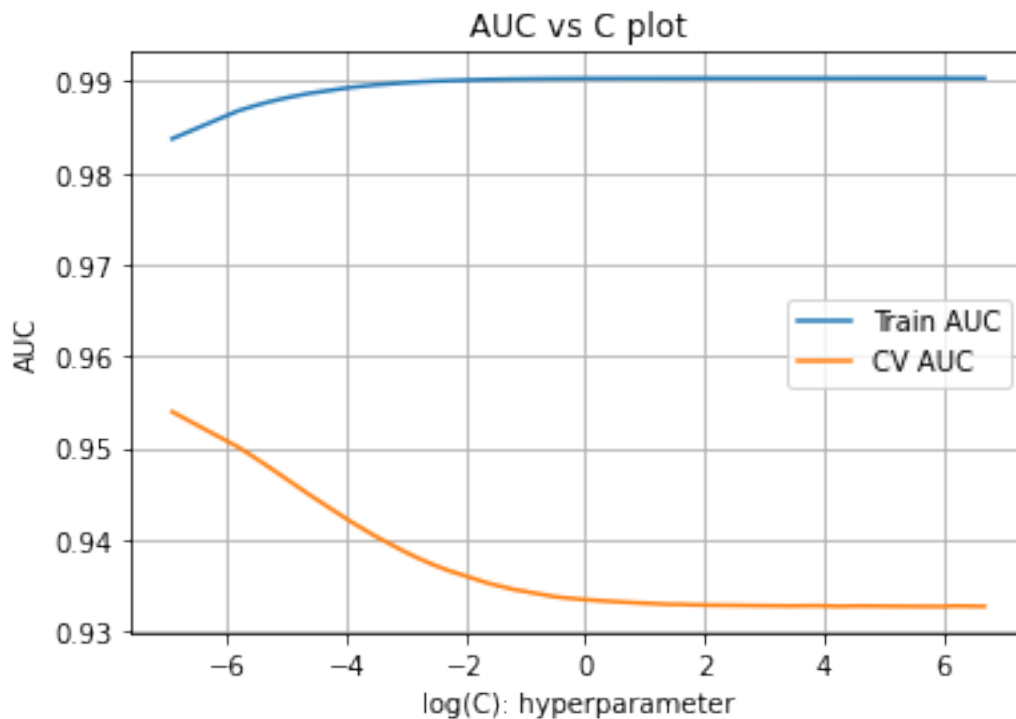
```



```

# plt.scatter(param['n_neighbors'], cv_auc, label='cv AUC')
plt.grid()
plt.title('AUC vs C plot')
plt.xlabel("log(C): hyperparameter")
plt.ylabel("AUC")
plt.legend()
plt.show()
plt.clf()
plt.cla()
plt.close()

```



[70]: *#finding the best CV score first then using the one which is least distant then its AUC counter part to avoid using Dumb model.*

```

from scipy.signal import argrelextrema
import numpy as np
x = np.array(train_auc)
y = np.array(cv_auc)
local_max=()
#finding local maximas of CV
local_max_i=argrelextrema(y, np.greater)
l=list(i for i in np.nditer(local_max_i))

```

```

diff=x-y
# diff between CV and Test AUC at the local maxima
local_diff=list(diff[i] for i in l)
local_c=list(c[i] for i in l)
print(f'all local differences {local_diff}')
print(f'all local max C {local_c}')

for i in np.nditer(np.argmin(local_diff)):
    v=i
    break
print(f'best cv score to use = {y[l[v]]}')
best_c=c[l[v]]
print(f'best C to use = {c[l[v]]}')

```

```

all local differences [0.05755916536927985, 0.057588965404848924,
0.057576852450763716]
all local max C [50, 100, 500]
best cv score to use = 0.9327657288890322
best C to use = 50

```

```

[0]: from sklearn.linear_model import LogisticRegression

from sklearn.metrics import
    →accuracy_score, confusion_matrix, f1_score, precision_score, recall_score, roc_curve,
    →auc

lr = LogisticRegression(penalty='l2', C=best_c)
lr.fit(X_train, y_train)
y_pred_tr = lr.predict_proba(X_train)
y_pred_ts = lr.predict_proba(X_test)
y_pred_ts=y_pred_ts[:,1]
y_pred_tr = y_pred_tr[:,1]


```

```

[72]: train_fpr, train_tpr, tr_thresholds = roc_curve(y_train, y_pred_tr)
test_fpr, test_tpr, te_thresholds = roc_curve(y_test, y_pred_ts)

plt.plot(train_fpr, train_tpr, label="train AUC =" + str(auc(train_fpr,
    →train_tpr)))
plt.plot(test_fpr, test_tpr, label="test AUC =" + str(auc(test_fpr, test_tpr)))
plt.xlabel("fpr")
plt.ylabel("tpr")
plt.title('ROC curve for ' + str(best_c) + ' as C')
plt.legend()
plt.grid()
plt.show()

```



[73]: *# This section of code where ever implemented is taken from sample kNN python_*
→notebook

```
def find_best_threshold(threshold, fpr, tpr):
    t = threshold[np.argmax(tpr*(1-fpr))]
    # (tpr*(1-fpr)) will be maximum if your fpr is very low and tpr is very_
    →high
    print("the maximum value of tpr*(1-fpr)", max(tpr*(1-fpr)), "for_
    →threshold", np.round(t,3))
    return t

def predict_with_best_t(proba, threshold):
    predictions = []
    for i in proba:
        if i>=threshold:
            predictions.append(1)
        else:
            predictions.append(0)
    return predictions

print('test')
best_ts_thres = find_best_threshold(te_thresholds, test_fpr, test_tpr)
```

```
print('train')
best_tr_thres = find_best_threshold(tr_thresholds, train_fpr, train_tpr)
```

test

the maximum value of $tpr \cdot (1 - fpr)$ 0.7480125959740773 for threshold 0.928


train

the maximum value of $tpr \cdot (1 - fpr)$ 0.8858638796688609 for threshold 0.797

```
[74]: print('train Confusion Matrix')
cm2 = pd.DataFrame(confusion_matrix(y_train, predict_with_best_t(y_pred_tr,
    ↳best_tr_thres)), range(2),range(2))
sns.heatmap(cm2, annot=True,fmt='g')
```

train Confusion Matrix


```
[74]: <matplotlib.axes._subplots.AxesSubplot at 0x7f3a4a7a3518>
```



```
[75]: print('Test Confusion Matrix')
cm2 = pd.DataFrame(confusion_matrix(y_test, predict_with_best_t(y_pred_ts,
    ↳best_ts_thres)), range(2),range(2))
sns.heatmap(cm2, annot=True,fmt='g')
```

Test Confusion Matrix

[75]: <matplotlib.axes._subplots.AxesSubplot at 0x7f3a4a4cd860>



```
[76]: acc=accuracy_score(y_test, predict_with_best_t(y_pred_ts, best_ts_thres))*100
ps=precision_score(y_test, predict_with_best_t(y_pred_ts, best_ts_thres))*100
rc=recall_score(y_test, predict_with_best_t(y_pred_ts, best_ts_thres))*100
f1=f1_score(y_test, predict_with_best_t(y_pred_ts, best_ts_thres))*100

print("Accuracy on test set: %.2f%%"%(acc))
print("Precision on test set: %.2f%%"%(ps))
print("recall score on test set: %.2f%%"%(rc))
print("f1 score on test set: %.2f%%"%(f1))
```

Accuracy on test set: 86.84%
Precision on test set: 96.99%
recall score on test set: 87.01%
f1 score on test set: 91.73%

```
[0]: count=0
value=[]
for i in lr.coef_.reshape(-1,1):
    count+=1
    value.extend(i)

x=vectorizer.get_feature_names()

log_features= pd.DataFrame({'feature_name':x,'value':value})
```

```

[78]: epsilon = 0.00002
      X_train.data = X_train.data +(epsilon,)
      X_train

[78]: <61441x5000 sparse matrix of type '<class 'numpy.float64'>'
      with 2079039 stored elements in Compressed Sparse Row format>

[0]: from sklearn.linear_model import LogisticRegression

      from sklearn.metrics import
      →accuracy_score,confusion_matrix,f1_score,precision_score,recall_score,roc_curve,
      →auc

      xlr = LogisticRegression(penalty='l2',C=best_c)
      xlr.fit(X_train,y_train)
      y_pred_tr = xlr.predict_proba(X_train)
      y_pred_ts = xlr.predict_proba(X_test)
      y_pred_ts=y_pred_ts[:,1]
      y_pred_tr = y_pred_tr[:,1]


[0]: w1=lr.coef_+0.000001
      w2=xlr.coef_+0.000001

[0]: p=abs((w1[0]-w2[0])/w1[0]*100)
      p=list(i for i in p)

[82]: change=pd.DataFrame({'change_per':(p), 'feature':x}, index=range(count))
      change.sort_values(by='change_per',ascending=False, inplace=True)
      # change=change[change['change_per']>0]
      plt.plot(range(change.shape[0]),np.log(change['change_per']))

[82]: [<matplotlib.lines.Line2D at 0x7f3a4b6706d8>]

```



```
[0]: '''
Algorithm to find elbow of a graph is taken from the following questionnaire on
↳Stackoverflow
#####
https://stackoverflow.com/questions/2018178/
↳finding-the-best-trade-off-point-on-a-curve
#####
'''
def elbow_finder(x_values):
    import numpy as np
    import numpy.matlib
    nPoints = len(x_values)
    allCoord = np.vstack((range(nPoints), x_values)).T
    np.array([range(nPoints), x_values])
    firstPoint = allCoord[0]
    lineVec = allCoord[-1] - allCoord[0]
    lineVecNorm = lineVec / np.sqrt(np.sum(lineVec**2))
    vecFromFirst = allCoord - firstPoint
    scalarProduct = np.sum(vecFromFirst * np.matlib.repmat(lineVecNorm, nPoints, 1), axis=1)
    vecFromFirstParallel = np.outer(scalarProduct, lineVecNorm)
    vecToLine = vecFromFirst - vecFromFirstParallel
    distToLine = np.sqrt(np.sum(vecToLine ** 2, axis=1))
    return np.argmax(distToLine)
```

```
[84]: x=elbow_finder(change['change_per'])
print(f'all the index before {x} is multicollinear as this is the elbow point')
```

all the index before 18 is multicollinear as this is the elbow point

```
[85]: change.iloc[:x]
```

```
[85]:      change_per      feature
342    16267.537796    better deal
2710   1878.043663      mix not
1577    574.235603      flavour
699     353.429829    chicken soup
3829    240.870293      section
1411    211.093968      fairly
3985    122.317819      smart
3926    115.609835    side effects
222      89.641187      auto ship
1303     76.031258      enough
1945     75.519137      gummies
2038     66.389935      hesitate
843      65.639399    comparable
2975     60.989554      not seen
1943     57.609354      gummi
4373     57.290151      tea also
1914     54.456117      grey tea
2132     45.787814    immediately
```

7.2.3 [5.2.3] Feature Importance on TFIDE, SET 2

[5.2.3.1] Top 10 important features of positive class from SET 2

```
[86]: log_features.sort_values(by = ['value'], ascending=False).head(10)
```

```
[86]:      feature_name      value
1872          great    1.043120
3797    say enough    1.020819
3698    right size    0.958855
1072    delicious    0.798143
1348    every penny    0.698103
3695    right amount    0.648381
2969    not regret    0.639794
325          best    0.592362
1791          good    0.590831
2833    no bitter    0.566061
```

[5.2.3.2] Top 10 important features of negative class from SET 2

```
[87]: log_features.sort_values(by = ['value'], ascending=True).head(10)
```



```
[87]:      feature_name      value
      2867          not -0.564613
      403         bob red -0.403381
     1145    disappointed -0.392602
     4906          worst -0.346650
     3003     not worth -0.332548
     3268         pg tips -0.310997
     4411     terrible -0.285917
     235         awful -0.278173
     374         bitter -0.261726
     4651    united states -0.257312
```

7.3 [5.3] Logistic Regression on AVG W2V, SET 3

7.3.1 [5.3.1] Applying Logistic Regression with L1 regularization on AVG W2V SET 3

```
[164]: from sklearn.model_selection import train_test_split
      from sklearn.preprocessing import StandardScaler

      X_train, X_test, y_train, y_test = train_test_split(sent_vectors, final['Score'].
      ↪ values, test_size=0.3, random_state=0)

      ss = StandardScaler(with_mean = False)
      X_train = ss.fit_transform(X_train)
      X_test = ss.transform(X_test)

      print(X_train.shape)
      print(X_test.shape)
```

```
(61441, 50)
```

```
(26332, 50)
```

```
[165]: param = {'C':c}
      print(param)


      from sklearn.model_selection import GridSearchCV
      from sklearn.linear_model import LogisticRegression
      l1= LogisticRegression(penalty='l1', solver='liblinear')
      temp_gscv=
      ↪ GridSearchCV(l1,param,cv=5,verbose=5,n_jobs=-1,scoring='roc_auc',return_train_score=True)
      temp_gscv.fit(X_train,y_train)
      temp_gs = pd.DataFrame.from_dict(temp_gscv.cv_results_)
```

```
{'C': [0.001, 0.002, 0.003, 0.004, 0.005, 0.007, 0.008, 0.01, 0.02, 0.03, 0.05,
0.06, 0.08, 0.1, 0.2, 0.3, 0.4, 0.5, 0.6, 0.8, 1, 2, 3, 4, 6, 7, 8, 10, 20, 40,
50, 80, 100, 200, 300, 400, 500, 600, 700, 800]}
```

Fitting 5 folds for each of 40 candidates, totalling 200 fits

```
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 2 concurrent workers.  
[Parallel(n_jobs=-1)]: Done 14 tasks      | elapsed: 6.1s  
[Parallel(n_jobs=-1)]: Done 68 tasks      | elapsed: 1.8min  
[Parallel(n_jobs=-1)]: Done 158 tasks     | elapsed: 6.9min  
[Parallel(n_jobs=-1)]: Done 200 out of 200 | elapsed: 9.3min finished
```

```
[166]: train_auc= temp_gs['mean_train_score']  
cv_auc= temp_gs['mean_test_score']  
plt.plot(np.log(c),train_auc,label='Train AUC')  
plt.plot(np.log(c),cv_auc ,label='CV AUC')  
  
# plt.scatter(param['n_neighbors'],train_auc,label='Train AUC')  
# plt.scatter(param['n_neighbors'], cv_auc, label='cv AUC')  
plt.grid()  
plt.title('AUC vs C plot')  
plt.xlabel("log(C): hyperparameter")  
plt.ylabel("AUC")  
plt.legend()  
plt.show()  
plt.clf()  
plt.cla()  
plt.close()
```



[167]: *#finding the best CV score first then using the one which is least distant then*
→its AUC counter part to avoid using Dumb model.

```
from scipy.signal import argrelextrema
import numpy as np
x = np.array(train_auc)
y = np.array(cv_auc)
local_max=()
#finding local maximas of CV
local_max_i=argrelextrema(y, np.greater)
l=list(i for i in np.nditer(local_max_i))

diff=x-y
# diff between CV and Test AUC at the local maxima
local_diff=list(diff[i] for i in l)
local_c=list(c[i] for i in l)
print(f'all local differences {local_diff}')
print(f'all local max C {local_c}')

for i in np.nditer(np.argmin(local_diff)):
    v=i
    break
print(f'best cv score to use = {y[l[v]]}')
best_c=c[l[v]]
print(f'best C to use = {c[l[v]]}')
```

all local differences [0.000821613809953492, 0.0008362096679768172,
0.0008423100624019852, 0.0008422906014544651, 0.0008433967481119575,
0.0008431618556553611, 0.0008425029079874857, 0.0008398848453644581]
all local max C [0.08, 0.8, 3, 7, 40, 100, 300, 500]
best cv score to use = 0.9095973270895991
best C to use = 0.08


```
[0]: from sklearn.linear_model import LogisticRegression

from sklearn.metrics import accuracy_score, confusion_matrix, f1_score, precision_score, recall_score, roc_curve, auc

lr = LogisticRegression(penalty='l1',C=best_c, solver='liblinear')
lr.fit(X_train,y_train)
y_pred_tr = lr.predict_proba(X_train)
y_pred_ts = lr.predict_proba(X_test)
y_pred_ts=y_pred_ts[:,1]
y_pred_tr = y_pred_tr[:,1]
```

```
[169]: train_fpr, train_tpr, tr_thresholds = roc_curve(y_train, y_pred_tr)
test_fpr, test_tpr, te_thresholds = roc_curve(y_test, y_pred_ts)

plt.plot(train_fpr, train_tpr, label="train AUC =" + str(auc(train_fpr,
→train_tpr)))
plt.plot(test_fpr, test_tpr, label="test AUC =" + str(auc(test_fpr, test_tpr)))
plt.xlabel("fpr")
plt.ylabel("tpr")
plt.title('ROC curve for ' + str (best_c) + 'as C')
plt.legend()
plt.grid()
plt.show()
```



```
[170]: # This section of code where ever implemented is taken from sample kNN python
→notebook

def find_best_threshold(threshould, fpr, tpr):
    t = threshould[np.argmax(tpr*(1-fpr))]
    # (tpr*(1-fpr)) will be maximum if your fpr is very low and tpr is very
    →high
    print("the maximum value of tpr*(1-fpr)", max(tpr*(1-fpr)), "for
    →threshold", np.round(t,3))
    return t
```

```
def predict_with_best_t(proba, threshold):
    predictions = []
    for i in proba:
        if i>=threshold:
            predictions.append(1)
        else:
            predictions.append(0)
    return predictions

print('test')
best_ts_thres = find_best_threshold(te_thresholds, test_fpr, test_tpr)


print('train')
best_tr_thres = find_best_threshold(tr_thresholds, train_fpr, train_tpr)
```

test
the maximum value of $tpr*(1-fpr)$ 0.6901199081166209 for threshold 0.803
train
the maximum value of $tpr*(1-fpr)$ 0.6903026623801541 for threshold 0.829

```
[171]: print('train Confusion Matrix')
cm2 = pd.DataFrame(confusion_matrix(y_train, predict_with_best_t(y_pred_tr,
→best_tr_thres)), range(2),range(2))
sns.heatmap(cm2, annot=True,fmt='g')
```

train Confusion Matrix


```
[171]: <matplotlib.axes._subplots.AxesSubplot at 0x7f3a4b0fd828>
```



```
[172]: print('Test Confusion Matrix')
cm2 = pd.DataFrame(confusion_matrix(y_test, predict_with_best_t(y_pred_ts,
↪best_ts_thres)), range(2),range(2))
sns.heatmap(cm2, annot=True,fmt='g')
```

Test Confusion Matrix

```
[172]: <matplotlib.axes._subplots.AxesSubplot at 0x7f3a4ba6acc0>
```



```
[173]: acc=accuracy_score(y_test, predict_with_best_t(y_pred_ts, best_ts_thres))*100
ps=precision_score(y_test, predict_with_best_t(y_pred_ts, best_ts_thres))*100
rc=recall_score(y_test, predict_with_best_t(y_pred_ts, best_ts_thres))*100
f1=f1_score(y_test, predict_with_best_t(y_pred_ts, best_ts_thres))*100

print("Accuracy on test set: %0.2f%%"%(acc))
print("Precision on test set: %0.2f%%"%(ps))
print("recall score on test set: %0.2f%%"%(rc))
print("f1 score on test set: %0.2f%%"%(f1))
```

Accuracy on test set: 84.22%
Precision on test set: 95.96%
recall score on test set: 84.75%
f1 score on test set: 90.01%

```
[174]: non_zero_features=np.count_nonzero(lr.coef_)
print(f"Number of features with non-zero weights :{non_zero_features}")
print(f"Number of features with zero weights :{5000-non_zero_features}")
```

Number of features with non-zero weights :48
Number of features with zero weights :4952

7.3.2 [5.3.2] Applying Logistic Regression with L2 regularization on AVG W2V, SET 3

```
[51]: #generating random alpha values between 10^-5 to 10^5
from numpy import random
c=[]
#removing 0 as value from alpha
def remove_zero(a):
    while True:
        if 0 in a:
            a.pop(0)
            continue
        break
#removing dupliacte
def remove_dup(a):
    for i,j in enumerate(a):
        if i!=(len(a)-1):
            if j == a[i+1]:
                a.pop(i)

#generating alpha values one at a time
def generator(a):
    x,y=random.randint(-3,3),random.randint(0,9)
    z=round(y*10**x,abs(x))
    a.extend([z])
    return a.sort()

while len(c)!=40:
    generator(c)
    remove_dup(c)
    remove_zero(c)

# c=[1000,100,10,1,0.1,0.01,0.001,0.0001]
print(c)
print(len(c))
```

[0.001, 0.002, 0.003, 0.005, 0.006, 0.007, 0.008, 0.03, 0.04, 0.08, 0.1, 0.2,
0.3, 0.4, 0.5, 0.6, 0.7, 0.8, 1, 2, 3, 4, 5, 6, 8, 20, 30, 40, 50, 60, 70, 80,
100, 200, 300, 400, 500, 600, 700, 800]
40

```
[53]: from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler

X_train, X_test, y_train, y_test = train_test_split(sent_vectors, final['Score'],
→ values, test_size=0.3, random_state=0)

ss = StandardScaler(with_mean = False)
X_train = ss.fit_transform(X_train)
X_test = ss.transform(X_test)

print(X_train.shape)
print(X_test.shape)
```

(61441, 50)

(26332, 50)

```
[54]: param = {'C':c}
print(param)

from sklearn.model_selection import GridSearchCV
from sklearn.linear_model import LogisticRegression
l1= LogisticRegression(penalty='l2')
temp_gscv=
→ GridSearchCV(l1,param,cv=5,verbose=5,n_jobs=-1,scoring='roc_auc',return_train_score=True)
temp_gscv.fit(X_train,y_train)
temp_gs = pd.DataFrame.from_dict(temp_gscv.cv_results_)
```

```
{'C': [0.001, 0.002, 0.003, 0.005, 0.006, 0.007, 0.008, 0.03, 0.04, 0.08, 0.1,
0.2, 0.3, 0.4, 0.5, 0.6, 0.7, 0.8, 1, 2, 3, 4, 5, 6, 8, 20, 30, 40, 50, 60, 70,
80, 100, 200, 300, 400, 500, 600, 700, 800]}
```

Fitting 5 folds for each of 40 candidates, totalling 200 fits

```
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 2 concurrent workers.
```

```
[Parallel(n_jobs=-1)]: Done 14 tasks | elapsed: 5.4s
```

```
[Parallel(n_jobs=-1)]: Done 68 tasks | elapsed: 27.4s
```

```
[Parallel(n_jobs=-1)]: Done 158 tasks | elapsed: 1.1min
```

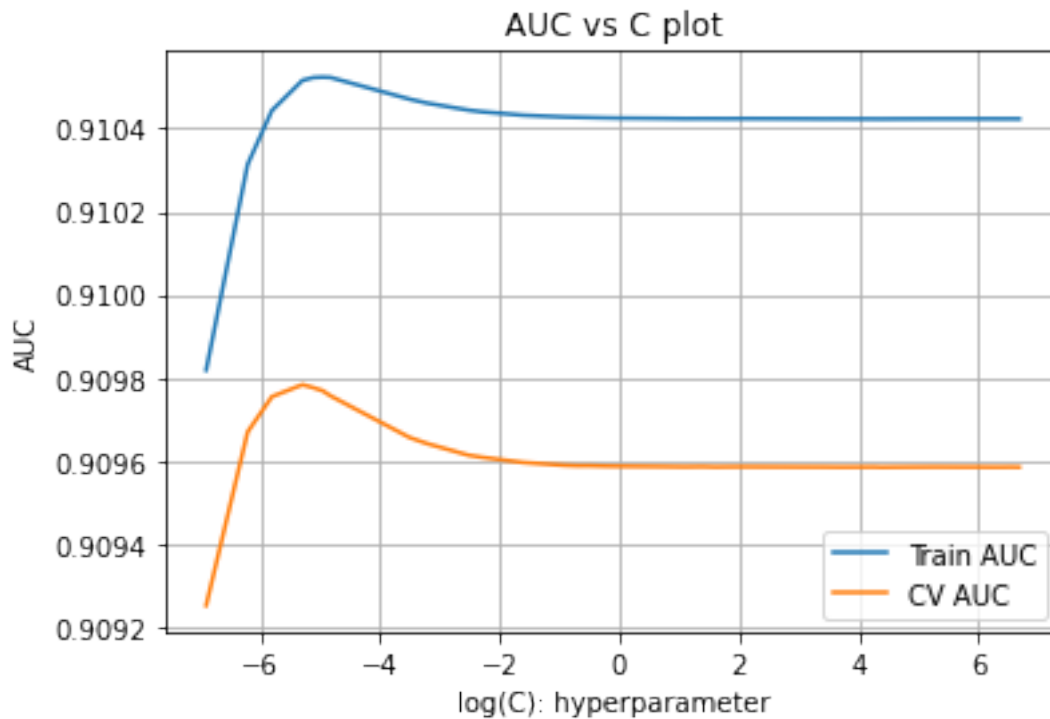
```
[Parallel(n_jobs=-1)]: Done 200 out of 200 | elapsed: 1.4min finished
```

```
[55]: train_auc= temp_gs['mean_train_score']
cv_auc= temp_gs['mean_test_score']
plt.plot(np.log(c),train_auc,label='Train AUC')
plt.plot(np.log(c),cv_auc ,label='CV AUC')

# plt.scatter(param['n_neighbors'],train_auc,label='Train AUC')
# plt.scatter(param['n_neighbors'], cv_auc, label='cv AUC')
plt.grid()
```



```
plt.title('AUC vs C plot')
plt.xlabel("log(C): hyperparameter")
plt.ylabel("AUC")
plt.legend()
plt.show()
plt.clf()
plt.cla()
plt.close()
```



[56]: *#finding the best CV score first then using the one which is least distant then_*
→ its AUC counter part to avoid using Dumb model.

```
from scipy.signal import argrelextrema
import numpy as np
x = np.array(train_auc)
y = np.array(cv_auc)
local_max=()
#finding local maximas of CV
local_max_i=argrelextrema(y, np.greater)
l=list(i for i in np.nditer(local_max_i))

diff=x-y
# diff between CV and Test AUC at the local maxima
```

```

local_diff=list(diff[i] for i in l)
local_c=list(c[i] for i in l)
print(f'all local differences {local_diff}')
print(f'all local max C {local_c}')

for i in np.nditer(np.argmin(local_diff)):
    v=i
    break
print(f'best cv score to use = {y[l[v]]}')
best_c=c[l[v]]
print(f'best C to use = {c[l[v]]}')

```

all local differences [0.0007325497170500572, 0.0008382568072425878,
0.0008378678855427424, 0.0008381982973316715, 0.000837914340929391,
0.0008380669820008402, 0.0008382166943246006, 0.0008382258839230161]
all local max C [0.005, 0.6, 4, 8, 70, 200, 400, 700]
best cv score to use = 0.90978336245962
best C to use = 0.005

```

[0]: from sklearn.linear_model import LogisticRegression


from sklearn.metrics import accuracy_score, confusion_matrix, f1_score, precision_score, recall_score, roc_curve, auc

lr = LogisticRegression(penalty='l2', C=best_c)
lr.fit(X_train, y_train)
y_pred_tr = lr.predict_proba(X_train)
y_pred_ts = lr.predict_proba(X_test)
y_pred_ts=y_pred_ts[:,1]
y_pred_tr = y_pred_tr[:,1]

[58]: train_fpr, train_tpr, tr_thresholds = roc_curve(y_train, y_pred_tr)
test_fpr, test_tpr, te_thresholds = roc_curve(y_test, y_pred_ts)

plt.plot(train_fpr, train_tpr, label="train AUC =" + str(auc(train_fpr, train_tpr)))
plt.plot(test_fpr, test_tpr, label="test AUC =" + str(auc(test_fpr, test_tpr)))
plt.xlabel("fpr")
plt.ylabel("tpr")
plt.title('ROC curve for ' + str(best_c) + 'as C')
plt.legend()
plt.grid()
plt.show()

```



[59]: *# This section of code where ever implemented is taken from sample kNN python notebook*

```
def find_best_threshold(threshold, fpr, tpr):
    t = threshold[np.argmax(tpr*(1-fpr))]
    # (tpr*(1-fpr)) will be maximum if your fpr is very low and tpr is very
    →high
    print("the maximum value of tpr*(1-fpr)", max(tpr*(1-fpr)), "for
    →threshold", np.round(t,3))
    return t

def predict_with_best_t(proba, threshold):
    predictions = []
    for i in proba:
        if i>=threshold:
            predictions.append(1)
        else:
            predictions.append(0)
    return predictions

print('test')
best_ts_thres = find_best_threshold(te_thresholds, test_fpr, test_tpr)
```

```
print('train')
best_tr_thres = find_best_threshold(tr_thresholds, train_fpr, train_tpr)
```

test

the maximum value of $tpr \cdot (1 - fpr)$ 0.6904011702791363 for threshold 0.804


train

the maximum value of $tpr \cdot (1 - fpr)$ 0.6902407972352401 for threshold 0.821

```
[60]: print('train Confusion Matrix')
cm2 = pd.DataFrame(confusion_matrix(y_train, predict_with_best_t(y_pred_tr,
    ↳best_tr_thres)), range(2),range(2))
sns.heatmap(cm2, annot=True,fmt='g')
```

train Confusion Matrix


```
[60]: <matplotlib.axes._subplots.AxesSubplot at 0x7f3a4eaec0b8>
```



```
[61]: print('Test Confusion Matrix')
cm2 = pd.DataFrame(confusion_matrix(y_test, predict_with_best_t(y_pred_ts,
    ↳best_ts_thres)), range(2),range(2))
sns.heatmap(cm2, annot=True,fmt='g')
```

Test Confusion Matrix

[61]: <matplotlib.axes._subplots.AxesSubplot at 0x7f3a4993a4e0>



```
[62]: acc=accuracy_score(y_test, predict_with_best_t(y_pred_ts, best_ts_thres))*100
ps=precision_score(y_test, predict_with_best_t(y_pred_ts, best_ts_thres))*100
rc=recall_score(y_test, predict_with_best_t(y_pred_ts, best_ts_thres))*100
f1=f1_score(y_test, predict_with_best_t(y_pred_ts, best_ts_thres))*100

print("Accuracy on test set: %0.2f%%"%(acc))
print("Precision on test set: %0.2f%%"%(ps))
print("recall score on test set: %0.2f%%"%(rc))
print("f1 score on test set: %0.2f%%"%(f1))
```

Accuracy on test set: 84.18%
Precision on test set: 95.97%
recall score on test set: 84.69%
f1 score on test set: 89.98%

7.4 [5.4] Logistic Regression on TFIDF W2V, SET 4

7.4.1 [5.4.1] Applying Logistic Regression with L2 regularization on TFIDF W2V, SET 4

```
[161]: from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
```

```

X_train, X_test, y_train, y_test =  
    → train_test_split(tfidf_sent_vectors, final['Score'].values, test_size=0.
    → 3, random_state=0)

ss = StandardScaler(with_mean = False)
X_train = ss.fit_transform(X_train)
X_test = ss.transform(X_test)

print(X_train.shape)
print(X_test.shape)

```

(61441, 50)

(26332, 50)

```

[162]: param = {'C':c}
        print(param)

        from sklearn.model_selection import GridSearchCV
        from sklearn.linear_model import LogisticRegression
        l1= LogisticRegression(penalty='l2')
        temp_gscv= 
            → GridSearchCV(l1,param,cv=5,verbose=5,n_jobs=-1,scoring='roc_auc',return_train_score=True)
        temp_gscv.fit(X_train,y_train)
        temp_gs = pd.DataFrame.from_dict(temp_gscv.cv_results_)

```

```

{'C': [0.001, 0.002, 0.003, 0.004, 0.005, 0.006, 0.008, 0.02, 0.04, 0.05, 0.06,
0.1, 0.2, 0.3, 0.4, 0.7, 0.8, 1, 2, 3, 4, 5, 6, 7, 8, 20, 30, 40, 50, 60, 70,
80, 100, 200, 300, 400, 500, 600, 700, 800]}

```

Fitting 5 folds for each of 40 candidates, totalling 200 fits

```

[Parallel(n_jobs=-1)]: Using backend LokyBackend with 2 concurrent workers.
[Parallel(n_jobs=-1)]: Done 14 tasks      | elapsed: 4.3s
[Parallel(n_jobs=-1)]: Done 68 tasks      | elapsed: 25.7s
[Parallel(n_jobs=-1)]: Done 158 tasks     | elapsed: 1.1min
[Parallel(n_jobs=-1)]: Done 200 out of 200 | elapsed: 1.4min finished

```


```

[151]: train_auc= temp_gs['mean_train_score']
        cv_auc= temp_gs['mean_test_score']
        plt.plot(np.log(c),train_auc,label='Train AUC')
        plt.plot(np.log(c),cv_auc ,label='CV AUC')

        # plt.scatter(param['n_neighbors'],train_auc,label='Train AUC')
        # plt.scatter(param['n_neighbors'], cv_auc, label='cv AUC')
        plt.grid()
        plt.title('AUC vs C plot')
        plt.xlabel("log(C): hyperparameter")
        plt.ylabel("AUC")

```

```
plt.legend()
plt.show()
plt.clf()
plt.cla()
plt.close()
```



[152]: *#finding the best CV score first then using the one which is least distant then, its AUC counter part to avoid using Dumb model.*

```
from scipy.signal import argrelextrema
import numpy as np
x = np.array(train_auc)
y = np.array(cv_auc)
local_max=()
#finding local maximas of CV
local_max_i=argrelextrema(y, np.greater)
l=list(i for i in np.nditer(local_max_i))

diff=x-y
# diff between CV and Test AUC at the local maxima
local_diff=list(diff[i] for i in l)
local_c=list(c[i] for i in l)
print(f'all local differences {local_diff}')
```

```

print(f'all local max C {local_c}')

for i in np.nditer(np.argmin(local_diff)):
    v=i
    break
print(f'best cv score to use = {y[l[v]]}')
best_c=c[l[v]]
print(f'best C to use = {c[l[v]]}')

```

```

all local differences [0.0701128003159106, 0.06581496103410067,
0.0655186103251224, 0.06523652497376142, 0.06518577171984141]
all local max C [3, 80, 200, 400, 700]
best cv score to use = 0.9348121320323994
best C to use = 700

```

```

[0]: from sklearn.linear_model import LogisticRegression

from sklearn.metrics import
    →accuracy_score,confusion_matrix,f1_score,precision_score,recall_score,roc_curve,
    →auc

```

```

lr = LogisticRegression(penalty='l2',C=best_c)
lr.fit(X_train,y_train)
y_pred_tr = lr.predict_proba(X_train)
y_pred_ts = lr.predict_proba(X_test)
y_pred_ts=y_pred_ts[:,1]
y_pred_tr = y_pred_tr[:,1]


```

```

[154]: train_fpr, train_tpr, tr_thresholds = roc_curve(y_train, y_pred_tr)
test_fpr, test_tpr, te_thresholds = roc_curve(y_test, y_pred_ts)

plt.plot(train_fpr, train_tpr, label="train AUC =" +str(auc(train_fpr,
    →train_tpr)))
plt.plot(test_fpr, test_tpr, label="test AUC =" +str(auc(test_fpr, test_tpr)))
plt.xlabel("fpr")
plt.ylabel("tpr")
plt.title('ROC curve for ' +str (best_c)+'as C')
plt.legend()
plt.grid()
plt.show()

```

[155]: *# This section of code where ever implemented is taken from sample kNN python_*
→notebook

```
def find_best_threshold(threshold, fpr, tpr):
    t = threshold[np.argmax(tpr*(1-fpr))]
    # (tpr*(1-fpr)) will be maximum if your fpr is very low and tpr is very_
    →high
    print("the maximum value of tpr*(1-fpr)", max(tpr*(1-fpr)), "for_
    →threshold", np.round(t,3))
    return t

def predict_with_best_t(proba, threshold):
    predictions = []
    for i in proba:
        if i>=threshold:
            predictions.append(1)
        else:
            predictions.append(0)
    return predictions

print('test')
best_ts_thres = find_best_threshold(te_thresholds, test_fpr, test_tpr)
```

```
print('train')
best_tr_thres = find_best_threshold(tr_thresholds, train_fpr, train_tpr)
```

test

the maximum value of $tpr \cdot (1 - fpr)$ 0.7364320906547198 for threshold 1.0


train

the maximum value of $tpr \cdot (1 - fpr)$ 0.9983310581298521 for threshold 0.889

```
[156]: print('train Confusion Matrix')
cm2 = pd.DataFrame(confusion_matrix(y_train, predict_with_best_t(y_pred_tr,
    ↳ best_tr_thres)), range(2), range(2))
sns.heatmap(cm2, annot=True, fmt='g')
```

train Confusion Matrix


```
[156]: <matplotlib.axes._subplots.AxesSubplot at 0x7f3a4b663a58>
```



```
[157]: print('Test Confusion Matrix')
cm2 = pd.DataFrame(confusion_matrix(y_test, predict_with_best_t(y_pred_ts,
    ↳ best_ts_thres)), range(2), range(2))
sns.heatmap(cm2, annot=True, fmt='g')
```

Test Confusion Matrix

[157]: <matplotlib.axes._subplots.AxesSubplot at 0x7f3a4d5e4240>



```
[158]: acc=accuracy_score(y_test, predict_with_best_t(y_pred_ts, best_ts_thres))*100
ps=precision_score(y_test, predict_with_best_t(y_pred_ts, best_ts_thres))*100
rc=recall_score(y_test, predict_with_best_t(y_pred_ts, best_ts_thres))*100
f1=f1_score(y_test, predict_with_best_t(y_pred_ts, best_ts_thres))*100

print("Accuracy on test set: %0.2f%%"%(acc))
print("Precision on test set: %0.2f%%"%(ps))
print("recall score on test set: %0.2f%%"%(rc))
print("f1 score on test set: %0.2f%%"%(f1))
```

Accuracy on test set: 85.46%
Precision on test set: 97.01%
recall score on test set: 85.29%
f1 score on test set: 90.77%

7.4.2 [5.4.2] Applying Logistic Regression with L1 regularization on TFIDF W2V, SET 4

```
[187]: from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler

X_train, X_test, y_train, y_test =
    →train_test_split(tfidf_sent_vectors, final['Score'].values, test_size=0.
    →3, random_state=0)
```

```

ss = StandardScaler(with_mean = False)
X_train = ss.fit_transform(X_train)
X_test = ss.transform(X_test)

print(X_train.shape)
print(X_test.shape)

```

```

(61441, 50)
(26332, 50)

```

```

[188]: param = {'C':c}
        print(param)

        from sklearn.model_selection import GridSearchCV
        from sklearn.linear_model import LogisticRegression
        l1= LogisticRegression(penalty='l1', solver='liblinear')
        temp_gscv=
        ↪GridSearchCV(l1,param,cv=5,verbose=5,n_jobs=-1,scoring='roc_auc',return_train_score=True)
        temp_gscv.fit(X_train,y_train)
        temp_gs = pd.DataFrame.from_dict(temp_gscv.cv_results_)

```

```

{'C': [0.001, 0.003, 0.004, 0.007, 0.02, 0.03, 0.04, 0.05, 0.06, 0.07, 0.08,
0.1, 0.2, 0.3, 0.4, 0.5, 0.6, 0.7, 0.8, 1, 2, 3, 4, 5, 6, 7, 8, 10, 30, 50, 70,
80, 100, 200, 300, 400, 500, 600, 700, 800]}

```

Fitting 5 folds for each of 40 candidates, totalling 200 fits

```

[Parallel(n_jobs=-1)]: Using backend LokyBackend with 2 concurrent workers.
[Parallel(n_jobs=-1)]: Done 14 tasks      | elapsed:    9.9s
[Parallel(n_jobs=-1)]: Done 68 tasks      | elapsed:   2.2min
[Parallel(n_jobs=-1)]: Done 158 tasks     | elapsed:   5.9min
[Parallel(n_jobs=-1)]: Done 200 out of 200 | elapsed:   7.7min finished

```


```

[189]: train_auc= temp_gs['mean_train_score']
        cv_auc= temp_gs['mean_test_score']
        plt.plot(np.log(c),train_auc,label='Train AUC')
        plt.plot(np.log(c),cv_auc ,label='CV AUC')

        # plt.scatter(param['n_neighbors'],train_auc,label='Train AUC')
        # plt.scatter(param['n_neighbors'], cv_auc, label='cv AUC')
        plt.grid()
        plt.title('AUC vs C plot')
        plt.xlabel("log(C): hyperparameter")
        plt.ylabel("AUC")
        plt.legend()
        plt.show()

```

```
plt.clf()
plt.cla()
plt.close()
```



[190]: *#finding the best CV score first then using the one which is least distant then*
→its AUC counter part to avoid using Dumb model.

```
from scipy.signal import argrelextrema
import numpy as np
x = np.array(train_auc)
y = np.array(cv_auc)
local_max=()
#finding local maximas of CV
local_max_i=argrelextrema(y, np.greater)
l=list(i for i in np.nditer(local_max_i))

diff=x-y
# diff between CV and Test AUC at the local maxima
local_diff=list(diff[i] for i in l)
local_c=list(c[i] for i in l)
print(f'all local differences {local_diff}')
print(f'all local max C {local_c}')
```

```

for i in np.nditer(np.argmin(local_diff)):
    v=i
    break
print(f'best cv score to use = {y[l[v]]}')
best_c=c[l[v]]
print(f'best C to use = {c[l[v]]}')

```

```

all local differences [0.0009249588816485632, 0.0009374255448282298,
0.0009368188325815652, 0.0009394946224473566, 0.0009434546224182938,
0.0009365700408060507, 0.0009350217442005704, 0.000941195484980617,
0.0009371853388489004]
all local max C [0.1, 0.7, 1, 4, 6, 8, 70, 300, 600]
best cv score to use = 0.8868188831531668
best C to use = 0.1

```

```

[0]: from sklearn.linear_model import LogisticRegression

from sklearn.metrics import
    →accuracy_score,confusion_matrix,f1_score,precision_score,recall_score,roc_curve,
    →auc

```

```

lr = LogisticRegression(penalty='l1',C=best_c, solver='liblinear')
lr.fit(X_train,y_train)
y_pred_tr = lr.predict_proba(X_train)
y_pred_ts = lr.predict_proba(X_test)
y_pred_ts=y_pred_ts[:,1]
y_pred_tr = y_pred_tr[:,1]


```

```

[192]: train_fpr, train_tpr, tr_thresholds = roc_curve(y_train, y_pred_tr)
test_fpr, test_tpr, te_thresholds = roc_curve(y_test, y_pred_ts)

plt.plot(train_fpr, train_tpr, label="train AUC =" +str(auc(train_fpr,
    →train_tpr)))
plt.plot(test_fpr, test_tpr, label="test AUC =" +str(auc(test_fpr, test_tpr)))
plt.xlabel("fpr")
plt.ylabel("tpr")
plt.title('ROC curve for ' +str (best_c)+'as C')
plt.legend()
plt.grid()
plt.show()

```



[193]: *# This section of code where ever implemented is taken from sample kNN python_*
→notebook

```
def find_best_threshold(threshold, fpr, tpr):
    t = threshold[np.argmax(tpr*(1-fpr))]
    # (tpr*(1-fpr)) will be maximum if your fpr is very low and tpr is very_
    →high
    print("the maximum value of tpr*(1-fpr)", max(tpr*(1-fpr)), "for_
    →threshold", np.round(t,3))
    return t

def predict_with_best_t(proba, threshold):
    predictions = []
    for i in proba:
        if i>=threshold:
            predictions.append(1)
        else:
            predictions.append(0)
    return predictions

print('test')
best_ts_thres = find_best_threshold(te_thresholds, test_fpr, test_tpr)
```


```
print('train')
best_tr_thres = find_best_threshold(tr_thresholds, train_fpr, train_tpr)
```

test
the maximum value of $tpr \cdot (1 - fpr)$ 0.6438650205123325 for threshold 0.821
train
the maximum value of $tpr \cdot (1 - fpr)$ 0.652366733483056 for threshold 0.823

```
[194]: print('train Confusion Matrix')
cm2 = pd.DataFrame(confusion_matrix(y_train, predict_with_best_t(y_pred_tr,
    ↳ best_tr_thres)), range(2), range(2))
sns.heatmap(cm2, annot=True, fmt='g')
```

train Confusion Matrix


```
[194]: <matplotlib.axes._subplots.AxesSubplot at 0x7f3a4b122160>
```



```
[195]: print('Test Confusion Matrix')
cm2 = pd.DataFrame(confusion_matrix(y_test, predict_with_best_t(y_pred_ts,
    ↳ best_ts_thres)), range(2), range(2))
sns.heatmap(cm2, annot=True, fmt='g')
```

Test Confusion Matrix

[195]: <matplotlib.axes._subplots.AxesSubplot at 0x7f3a51b75390>



```
[196]: acc=accuracy_score(y_test, predict_with_best_t(y_pred_ts, best_ts_thres))*100
ps=precision_score(y_test, predict_with_best_t(y_pred_ts, best_ts_thres))*100
rc=recall_score(y_test, predict_with_best_t(y_pred_ts, best_ts_thres))*100
f1=f1_score(y_test, predict_with_best_t(y_pred_ts, best_ts_thres))*100

print("Accuracy on test set: %0.2f%%"%(acc))
print("Precision on test set: %0.2f%%"%(ps))
print("recall score on test set: %0.2f%%"%(rc))
print("f1 score on test set: %0.2f%%"%(f1))
```

Accuracy on test set: 80.88%
Precision on test set: 95.33%
recall score on test set: 81.19%
f1 score on test set: 87.69%

```
[198]: non_zero_features=np.count_nonzero(lr.coef_)
print(f"Number of features with non-zero weights :{non_zero_features}")
```

Number of features with non-zero weights :49

8 [6] Conclusions

```
[200]: from prettytable import PrettyTable
x = PrettyTable()
x.field_names = ["S.NO.", "MODEL", "C value", "Penalty", "Test AUC", "Precision_↵Score"]
x.add_row(["1", "BOW", "0.02", "L1", "0.9500", "97.52%"])
x.add_row(["", "BOW", "7", "L2", "0.9255", "96.71%"])
x.add_row(["2", "TFIDF", "0.02", "L1", "0.9583", "97.85%"])
x.add_row(["", "TFIDF", "50", "L2", "0.9393", "96.99%"])
x.add_row(["3", "AVG W2V", "0.08", "L1", "0.9074", "95.96%"])
x.add_row(["", "AVG W2V", "0.05", "L2", "0.9075", "95.97%"])
x.add_row(["4", "TFIDF W2V", "0.001", "L1", "0.8831", "95.33%"])
x.add_row(["", "TFIDF W2V", "700", "L2", "0.9303", "97.01%"])
print(x)
```

S.NO.	MODEL	C value	Penalty	Test AUC	Precision Score
1	BOW	0.02	L1	0.9500	97.52%
	BOW	7	L2	0.9255	96.71%
2	TFIDF	0.02	L1	0.9583	97.85%
	TFIDF	50	L2	0.9393	96.99%
3	AVG W2V	0.08	L1	0.9074	95.96%
	AVG W2V	0.05	L2	0.9075	95.97%
4	TFIDF W2V	0.001	L1	0.8831	95.33%
	TFIDF W2V	700	L2	0.9303	97.01%