

Amazon Fine Food Reviews Analysis

Data Source: <https://www.kaggle.com/snap/amazon-fine-food-reviews> (<https://www.kaggle.com/snap/amazon-fine-food-reviews>)

EDA: <https://nycdatascience.com/blog/student-works/amazon-fine-foods-visualization/>
(<https://nycdatascience.com/blog/student-works/amazon-fine-foods-visualization/>)

The Amazon Fine Food Reviews dataset consists of reviews of fine foods from Amazon.

Number of reviews: 568,454

Number of users: 256,059

Number of products: 74,258

Timespan: Oct 1999 - Oct 2012

Number of Attributes/Columns in data: 10

Attribute Information:

1. Id
2. ProductId - unique identifier for the product
3. UserId - unique identifier for the user
4. ProfileName
5. HelpfulnessNumerator - number of users who found the review helpful
6. HelpfulnessDenominator - number of users who indicated whether they found the review helpful or not
7. Score - rating between 1 and 5
8. Time - timestamp for the review
9. Summary - brief summary of the review
10. Text - text of the review

Objective:

Given a review, determine whether the review is positive (rating of 4 or 5) or negative (rating of 1 or 2).

[Q] How to determine if a review is positive or negative?

[Ans] We could use Score/Rating. A rating of 4 or 5 can be considered as a positive review. A rating of 1 or 2 can be considered as negative one. A review of rating 3 is considered neutral and such reviews are ignored from our analysis. This is an approximate and proxy way of determining the polarity (positivity/negativity) of a review.

[1]. Reading Data

[1.1] Loading the data

The dataset is available in two forms

1. .csv file
2. SQLite Database

In order to load the data, We have used the SQLITE dataset as it is easier to query the data and visualise the data efficiently.

Here as we only want to get the global sentiment of the recommendations (positive or negative), we will purposefully ignore all Scores equal to 3. If the score is above 3, then the recommendation will be set to "positive". Otherwise, it will be set to "negative".

In [0]:

```
1 %matplotlib inline
2 import warnings
3 warnings.filterwarnings("ignore")
4
5
6 import sqlite3
7 import pandas as pd
8 import numpy as np
9 import nltk
10 import string
11 import matplotlib.pyplot as plt
12 import seaborn as sns
13 from sklearn.feature_extraction.text import TfidfTransformer
14 from sklearn.feature_extraction.text import TfidfVectorizer
15 from sklearn.model_selection import TimeSeriesSplit
16
17 from sklearn.feature_extraction.text import CountVectorizer
18 from sklearn.metrics import confusion_matrix
19 from sklearn import metrics
20 from sklearn.metrics import roc_curve, auc
21 from nltk.stem.porter import PorterStemmer
22
23 import re
24 # Tutorial about Python regular expressions: https://pymotw.com/2/re/
25 import string
26 from nltk.corpus import stopwords
27 from nltk.stem import PorterStemmer
28 from nltk.stem.wordnet import WordNetLemmatizer
29
30 from gensim.models import Word2Vec
31 from gensim.models import KeyedVectors
32 import pickle
33
34 from tqdm import tqdm
35 import os
```

In [2]:

```
1 from google.colab import drive
2 drive.mount('/content/drive')
```

Go to this URL in a browser: https://accounts.google.com/o/oauth2/auth?client_id=947318989803-6bn6qk8qdgf4n4g3pfee6491hc0brc4i.apps.googleusercontent.com&redirect_uri=urn%3Aietf%3Awg%3Aoauth%3A2.0%3Aoob&scope=email%20https%3A%2F%2Fwww.googleapis.com%2Fauth%2Fdocs.test%20https%3A%2F%2Fwww.googleapis.com%2Fauth%2Fdrive%20https%3A%2F%2Fwww.googleapis.com%2Fauth%2Fdrive.photos.readonly%20https%3A%2F%2Fwww.googleapis.com%2Fauth%2Fpeopleapi.readonly&response_type=code (https://accounts.google.com/o/oauth2/auth?client_id=947318989803-6bn6qk8qdgf4n4g3pfee6491hc0brc4i.apps.googleusercontent.com&redirect_uri=urn%3Aietf%3Awg%3Aoauth%3A2.0%3Aoob&scope=email%20https%3A%2F%2Fwww.googleapis.com%2Fauth%2Fdocs.test%20https%3A%2F%2Fwww.googleapis.com%2Fauth%2Fdrive%20https%3A%2F%2Fwww.googleapis.com%2Fauth%2Fdrive.photos.readonly%20https%3A%2F%2Fwww.googleapis.com%2Fauth%2Fpeopleapi.readonly&response_type=code)

Enter your authorization code:

.....

Mounted at /content/drive

In [34]:

```

1  # using SQLite Table to read data.
2  con = sqlite3.connect('drive/My Drive/FFRDB/database.sqlite')
3
4  # filtering only positive and negative reviews i.e.
5  # not taking into consideration those reviews with Score=3
6  # SELECT * FROM Reviews WHERE Score != 3 LIMIT 500000, will give top 500000 data point
7  # you can change the number to any other number based on your computing power
8
9  # filtered_data = pd.read_sql_query(""" SELECT * FROM Reviews WHERE Score != 3 LIMIT 5
10 # for tsne assignment you can take 5k data points
11
12 filtered_data = pd.read_sql_query(""" SELECT * FROM Reviews WHERE Score != 3""", con)
13
14 # Give reviews with Score>3 a positive rating(1), and reviews with a score<3 a negativ
15 def partition(x):
16     if x < 3:
17         return 0
18     return 1
19
20 #changing reviews with score less than 3 to be positive and vice-versa
21 actualScore = filtered_data['Score']
22 positiveNegative = actualScore.map(partition)
23 filtered_data['Score'] = positiveNegative
24 print("Number of data points in our data", filtered_data.shape)
25 filtered_data.head(3)

```

Number of data points in our data (525814, 10)

Out[34]:

	Id	ProductId	UserId	ProfileName	HelpfulnessNumerator	HelpfulnessDenomr
0	1	B001E4KFG0	A3SGXH7AUHU8GW	delmartian	1	
1	2	B00813GRG4	A1D87F6ZCVE5NK	dll pa	0	
2	3	B000LQOCH0	ABXLMWJIXXAIN	Natalia Corres "Natalia Corres"	1	

In [0]:

```
1 display = pd.read_sql_query("""
2 SELECT UserId, ProductId, ProfileName, Time, Score, Text, COUNT(*)
3 FROM Reviews
4 GROUP BY UserId
5 HAVING COUNT(*)>1
6 """, con)
```

In [36]:

```
1 print(display.shape)
2 display.head()
```

(80668, 7)

Out[36]:

	UserId	ProductId	ProfileName	Time	Score	Text	COUNT(*)
0	#oc-R115TNMSPFT9I7	B007Y59HVM	Breyton	1331510400	2	Overall its just OK when considering the price...	2
1	#oc-R11D9D7SHXIJ9	B005HG9ET0	Louis E. Emory "hoppy"	1342396800	5	My wife has recurring extreme muscle spasms, u...	3
2	#oc-R11DNU2NBKQ23Z	B007Y59HVM	Kim Cieszykowski	1348531200	1	This coffee is horrible and unfortunately not ...	2
3	#oc-R11O5J5ZVQE25C	B005HG9ET0	Penguin Chick	1346889600	5	This will be the bottle that you grab from the...	3
4	#oc-R12KPBODL2B5ZD	B007OSBE1U	Christopher P. Presta	1348617600	1	I didnt like this coffee. Instead of telling y...	2

In [37]:

```
1 display[display['UserId']=='AZY10LLTJ71NX']
```

Out[37]:

	UserId	ProductId	ProfileName	Time	Score	Text	COUNT(*)
80638	AZY10LLTJ71NX	B006P7E5ZI	undertheshrine "undertheshrine"	1334707200	5	I was recommended to try green tea extract to ...	

In [38]:

```
1 display['COUNT(*)'].sum()
```

Out[38]:

393063

[2] Exploratory Data Analysis

[2.1] Data Cleaning: Deduplication

It is observed (as shown in the table below) that the reviews data had many duplicate entries. Hence it was necessary to remove duplicates in order to get unbiased results for the analysis of the data. Following is an example:

In [39]:

```

1 display= pd.read_sql_query("""
2 SELECT *
3 FROM Reviews
4 WHERE Score != 3 AND UserId="AR5J8UI46CURR"
5 ORDER BY ProductID
6 """, con)
7 display.head()

```

Out[39]:

	Id	ProductId	UserId	ProfileName	HelpfulnessNumerator	HelpfulnessDenominator
0	78445	B000HDL1RQ	AR5J8UI46CURR	Geetha Krishnan	2	2
1	138317	B000HDOPYC	AR5J8UI46CURR	Geetha Krishnan	2	2
2	138277	B000HDOPYM	AR5J8UI46CURR	Geetha Krishnan	2	2
3	73791	B000HDOPZG	AR5J8UI46CURR	Geetha Krishnan	2	2
4	155049	B000PAQ75C	AR5J8UI46CURR	Geetha Krishnan	2	2

As it can be seen above that same user has multiple reviews with same values for HelpfulnessNumerator, HelpfulnessDenominator, Score, Time, Summary and Text and on doing analysis it was found that

ProductId=B000HDOPZG was Loacker Quadratini Vanilla Wafer Cookies, 8.82-Ounce Packages (Pack of 8)

ProductId=B000HDL1RQ was Loacker Quadratini Lemon Wafer Cookies, 8.82-Ounce Packages (Pack of 8) and so on

It was inferred after analysis that reviews with same parameters other than ProductId belonged to the same product just having different flavour or quantity. Hence in order to reduce redundancy it was decided to eliminate the rows having same parameters.

The method used for the same was that we first sort the data according to ProductId and then just keep the first similar product review and delete the others. for eg. in the above just the review for ProductId=B000HDL1RQ remains. This method ensures that there is only one representative for each product and deduplication without sorting would lead to possibility of different representatives still existing for the same product.

In [0]:

```
1 #Sorting data according to ProductId in ascending order
2
3 sorted_data=filtered_data.sort_values('ProductId', axis=0, ascending=True, inplace=False)
```

In [41]:

```
1 #Deduplication of entries
2 final=sorted_data.drop_duplicates(subset={"UserId","ProfileName","Time","Text"}, keep='first')
3
4 # sorting data based on time and taking first 100000 observations
5 final=final.sort_values('Time', axis=0, ascending=True, inplace=False, kind='quicksort')
6 final=final[:100000]
7 final.shape
```

Out[41]:

(100000, 10)

In [42]:

```
1 #Checking to see how much % of data still remains
2 (final['Id'].size*1.0)/(filtered_data['Id'].size*1.0)*100
```

Out[42]:

19.01813188694101

Observation:- It was also seen that in two rows given below the value of HelpfulnessNumerator is greater than HelpfulnessDenominator which is not practically possible hence these two rows too are removed from calculations

In [43]:

```

1 display= pd.read_sql_query("""
2 SELECT *
3 FROM Reviews
4 WHERE Score != 3 AND Id=44737 OR Id=64422
5 ORDER BY ProductID
6 """, con)
7
8 display.head()

```

Out[43]:

	Id	ProductId	UserId	ProfileName	HelpfulnessNumerator	HelpfulnessDenominator
0	64422	B000MIDROQ	A161DK06JJMCYF	J. E. Stephens "Jeanne"	3	
1	44737	B001EQ55RW	A2V0I904FH7ABY	Ram	3	

In [0]:

```
1 final=final[final.HelpfulnessNumerator<=final.HelpfulnessDenominator]
```

In [45]:

```

1 #Before starting the next phase of preprocessing Lets see the number of entries left
2 print(final.shape)
3
4 #How many positive and negative reviews are present in our dataset?
5 final['Score'].value_counts()

```

(99998, 10)

Out[45]:

```

1      87729
0      12269
Name: Score, dtype: int64

```

[3] Preprocessing

[3.1]. Preprocessing Review Text

Now that we have finished deduplication our data requires some preprocessing before we go on further with analysis and making the prediction model.

Hence in the Preprocessing phase we do the following in the order below:-

1. Begin by removing the html tags
2. Remove any punctuations or limited set of special characters like , or . or # etc.
3. Check if the word is made up of english letters and is not alpha-numeric
4. Check to see if the length of the word is greater than 2 (as it was researched that there is no adjective in 2-letters)
5. Convert the word to lowercase
6. Remove Stopwords
7. Finally Snowball Stemming the word (it was observed to be better than Porter Stemming)

After which we collect the words used to describe positive and negative reviews

In [46]:

```

1  # printing some random reviews
2  sent_0 = final['Text'].values[0]
3  print(sent_0)
4  print("="*50)
5
6  sent_1000 = final['Text'].values[1000]
7  print(sent_1000)
8  print("="*50)
9
10 sent_1500 = final['Text'].values[1500]
11 print(sent_1500)
12 print("="*50)
13
14 sent_4900 = final['Text'].values[4900]
15 print(sent_4900)
16 print("="*50)

```

this witty little book makes my son laugh at loud. i recite it in the car as we're driving along and he always can sing the refrain. he's learned about whales, India, drooping roses: i love all the new words this book introduce s and the silliness of it all. this is a classic book i am willing to bet my son will STILL be able to recite from memory when he is in college

=====

I finally ordered a couple products from this seller for myself(not as gift s) and I am really happy. This Jade Bonsai is really cool and it arrived fast and in perfect condition. It's in my living room and I get tons of compliments. It's already grown some too and the pot it came in is really nice, looks expensive! Much bigger than I thought it would be even. Thanks again!!

=====

I bought some of this tea when I was in Seattle and I have been dying to get more. It really is the best tea I have ever had. It is great hot or cold.

=====

I would prefer freshly made brown rice, but that takes a long time to make and isn't easy. This makes it convenient, and takes all the guess work out of making it. I generally have been buying frozen organic brown rice, but that takes up lots of freezer space. The fact that this is easy to store at room temperature is a big plus. I'll be buying more.

=====

In [47]:

```
1 # remove urls from text python: https://stackoverflow.com/a/40823105/4084039
2 sent_0 = re.sub(r"http\S+", "", sent_0)
3 sent_1000 = re.sub(r"http\S+", "", sent_1000)
4 sent_150 = re.sub(r"http\S+", "", sent_1500)
5 sent_4900 = re.sub(r"http\S+", "", sent_4900)
6
7 print(sent_0)
```

this witty little book makes my son laugh at loud. i recite it in the car as we're driving along and he always can sing the refrain. he's learned about whales, India, drooping roses: i love all the new words this book introduces and the silliness of it all. this is a classic book i am willing to bet my son will STILL be able to recite from memory when he is in college

In [48]:

```

1  # https://stackoverflow.com/questions/16206380/python-beautifulsoup-how-to-remove-all-
2  from bs4 import BeautifulSoup
3
4  soup = BeautifulSoup(sent_0, 'lxml')
5  text = soup.get_text()
6  print(text)
7  print("="*50)
8
9  soup = BeautifulSoup(sent_1000, 'lxml')
10 text = soup.get_text()
11 print(text)
12 print("="*50)
13
14 soup = BeautifulSoup(sent_1500, 'lxml')
15 text = soup.get_text()
16 print(text)
17 print("="*50)
18
19 soup = BeautifulSoup(sent_4900, 'lxml')
20 text = soup.get_text()
21 print(text)

```

this witty little book makes my son laugh at loud. i recite it in the car as we're driving along and he always can sing the refrain. he's learned about whales, India, drooping roses: i love all the new words this book introduce s and the silliness of it all. this is a classic book i am willing to bet my son will STILL be able to recite from memory when he is in college

=====

I finally ordered a couple products from this seller for myself(not as gift s) and I am really happy. This Jade Bonsai is really cool and it arrived fast and in perfect condition. It's in my living room and I get tons of compliments. It's already grown some too and the pot it came in is really nice, looks expensive! Much bigger than I thought it would be even. Thanks again!!

=====

I bought some of this tea when I was in Seattle and I have been dying to get more. It really is the best tea I have ever had. It is great hot or cold.

=====

I would prefer freshly made brown rice, but that takes a long time to make and isn't easy. This makes it convenient, and takes all the guess work out of making it. I generally have been buying frozen organic brown rice, but that takes up lots of freezer space. The fact that this is easy to store at room temperature is a big plus. I'll be buying more.

In [0]:

```

1 # https://stackoverflow.com/a/47091490/4084039
2 import re
3
4 def decontracted(phrase):
5     # specific
6     phrase = re.sub(r"won't", "will not", phrase)
7     phrase = re.sub(r"can't", "can not", phrase)
8
9     # general
10    phrase = re.sub(r"n't", " not", phrase)
11    phrase = re.sub(r"'re", " are", phrase)
12    phrase = re.sub(r"'s", " is", phrase)
13    phrase = re.sub(r"'d", " would", phrase)
14    phrase = re.sub(r"'ll", " will", phrase)
15    phrase = re.sub(r"'t", " not", phrase)
16    phrase = re.sub(r"'ve", " have", phrase)
17    phrase = re.sub(r"'m", " am", phrase)
18    return phrase

```

In [50]:

```

1 sent_1500 = decontracted(sent_1500)
2 print(sent_1500)
3 print("="*50)

```

I bought some of this tea when I was in Seattle and I have been dying to get more. It really is the best tea I have ever had. It is great hot or cold.
=====

In [51]:

```

1 #remove words with numbers python: https://stackoverflow.com/a/18082370/4084039
2 sent_0 = re.sub(r"\S*\d\S*", "", sent_0).strip()
3 print(sent_0)

```

this witty little book makes my son laugh at loud. i recite it in the car as we're driving along and he always can sing the refrain. he's learned about whales, India, drooping roses: i love all the new words this book introduces and the silliness of it all. this is a classic book i am willing to bet my son will STILL be able to recite from memory when he is in college

In [52]:

```

1 #remove spacial character: https://stackoverflow.com/a/5843547/4084039
2 sent_1500 = re.sub(r'[^A-Za-z0-9]+', ' ', sent_1500)
3 print(sent_1500)

```

I bought some of this tea when I was in Seattle and I have been dying to get more It really is the best tea I have ever had It is great hot or cold

In [0]:

```

1 # https://gist.github.com/sebleier/554280
2 # we are removing the words from the stop words list: 'no', 'nor', 'not'
3 # <br /><br /> ==> after the above steps, we are getting "br br"
4 # we are including them into stop words list
5 # instead of <br /> if we have <br/> these tags would have revmoved in the 1st step
6
7 stopwords= set(['br', 'the', 'i', 'me', 'my', 'myself', 'we', 'our', 'ours', 'ourself',
8                "you'll", "you'd", 'your', 'yours', 'yourself', 'yourselves', 'he', 'him',
9                'she', "she's", 'her', 'hers', 'herself', 'it', "it's", 'its', 'itself', 'his',
10               'theirs', 'themselves', 'what', 'which', 'who', 'whom', 'this', 'that', 'they',
11               'am', 'is', 'are', 'was', 'were', 'be', 'been', 'being', 'have', 'has', 'had',
12               'did', 'doing', 'a', 'an', 'the', 'and', 'but', 'if', 'or', 'because', 'as',
13               'at', 'by', 'for', 'with', 'about', 'against', 'between', 'into', 'through',
14               'above', 'below', 'to', 'from', 'up', 'down', 'in', 'out', 'on', 'off', 'over',
15               'then', 'once', 'here', 'there', 'when', 'where', 'why', 'how', 'all', 'any',
16               'most', 'other', 'some', 'such', 'only', 'own', 'same', 'so', 'than', 'too',
17               's', 't', 'can', 'will', 'just', 'don', "don't", 'should', "should've", 'n',
18               've', 'y', 'ain', 'aren', "aren't", 'couldn', "couldn't", 'didn', "didn't",
19               "hadn't", 'hasn', "hasn't", 'haven', "haven't", 'isn', "isn't", 'ma', 'mightn',
20               "mustn't", 'needn', "needn't", 'shan', "shan't", 'shouldn', "shouldn't", 'w',
21               'won', "won't", 'wouldn', "wouldn't"])

```

In [54]:

```

1 # Combining all the above stundents
2 from tqdm import tqdm
3 preprocessed_reviews = []
4 # tqdm is for printing the status bar
5 for sentence in tqdm(final['Text'].values):
6     sentence = re.sub(r"http\S+", "", sentence)
7     sentence = BeautifulSoup(sentence, 'lxml').get_text()
8     sentence = decontracted(sentence)
9     sentence = re.sub("\S*\d\S*", "", sentence).strip()
10    sentence = re.sub('[^A-Za-z]+', ' ', sentence)
11    # https://gist.github.com/sebleier/554280
12    sentence = ' '.join(e.lower() for e in sentence.split() if e.lower() not in stopwords)
13    preprocessed_reviews.append(sentence.strip())

```

100%|██████████| 99998/99998 [00:39<00:00, 2532.56it/s]

In [55]:

```
1 preprocessed_reviews[1500]
```

Out[55]:

'bought tea seattle dying get really best tea ever great hot cold'

[5] Assignment 4: Apply Naive Bayes

1. Apply Multinomial NaiveBayes on these feature sets

- **SET 1:** Review text, preprocessed one converted into vectors using (BOW)
- **SET 2:** Review text, preprocessed one converted into vectors using (TFIDF)

2. The hyper paramter tuning(find best Alpha)

- Find the best hyper parameter which will give the maximum [AUC](https://www.appliedaicourse.com/course/applied-ai-course-online/lessons/receiver-operating-characteristic-curve-roc-curve-and-auc-1/) value
- Consider a wide range of alpha values for hyperparameter tuning, start as low as 0.00001
- Find the best hyper paramter using k-fold cross validation or simple cross validation data
- Use gridsearch cv or randomsearch cv or you can also write your own for loops to do this task of hyperparameter tuning

3. Feature importance

- Find the top 10 features of positive class and top 10 features of negative class for both feature sets **Set 1** and **Set 2** using values of `feature_log_prob_` parameter of [MultinomialNB](https://scikit-learn.org/stable/modules/generated/sklearn.naive_bayes.MultinomialNB.html) and print their corresponding feature names

4. Feature engineering

- To increase the performance of your model, you can also experiment with with feature engineering like :
 - Taking length of reviews as another feature.
 - Considering some features from review summary as well.

5. Representation of results

- You need to plot the performance of model both on train data and cross validation data for each hyper parameter, like shown in the figure. Here on X-axis you will have alpha values, since they have a wide range, just to represent those alpha values on the graph, apply log function on those alpha values.
- Once after you found the best hyper parameter, you need to train your model with it, and find the AUC on test data and plot the ROC curve on both train and test.
- Along with plotting ROC curve, you need to print the [confusion matrix](https://www.appliedaicourse.com/course/applied-ai-course-online/lessons/confusion-matrix-tpr-fpr-fnr-tnr-1/) with predicted and original labels of test data points. Please visualize your confusion matrices using [seaborn heatmaps](https://seaborn.pydata.org/generated/seaborn.heatmap.html).

6. Conclusion

- You need to summarize the results at the end of the notebook, summarize it in the table format. To print out a table please refer to this prettytable library [link](http://zetcode.com/python/prettytable/)

Note: Data Leakage

1. There will be an issue of data-leakage if you vectorize the entire data and then split it into train/cv/test.
2. To avoid the issue of data-leakag, make sure to split your data first and then vectorize it.

3. While vectorizing your data, apply the method `fit_transform()` on you train data, and apply the method `transform()` on cv/test data.
4. For more details please go through this [link. \(https://soundcloud.com/applied-ai-course/leakage-bow-and-tfidf\)](https://soundcloud.com/applied-ai-course/leakage-bow-and-tfidf)

Applying Multinomial Naive Bayes

[5.1] Applying Naive Bayes on BOW, SET 1

In [118]:

```
1 from sklearn.preprocessing import StandardScaler
2 from sklearn.feature_extraction.text import CountVectorizer
3 from sklearn.model_selection import train_test_split
4
5 ## splitting with shuffle as false so time consistency is maintained as data is sorted
6
7 X_train_bow, X_test_bow, y_train_bow, y_test_bow = train_test_split(preprocessed_review, y,
8
9 BOW = CountVectorizer(ngram_range=(1,2), min_df=10, max_features=50000)
10 BOW.fit_transform(X_train_bow)
11 X_train_bow = BOW.transform(X_train_bow)
12 X_test_bow = BOW.transform(X_test_bow)
13
14 scaler = StandardScaler(with_mean = False)
15 X_train_bow = scaler.fit_transform(X_train_bow)
16 X_test_bow = scaler.transform(X_test_bow)
17
18 print(X_train_bow.shape)
19 print(X_test_bow.shape)
```

(69998, 39325)

(30000, 39325)

In [119]:

```
1 #generating random alpha values between 10^-5 to 10^5
2 from numpy import random
3 alpha=list()
4 alpha.extend(list(map(lambda x,y:y*10**x,list(range(-5,5)),list(random.randint(0,9,siz
5 alpha.extend(list(map(lambda x,y:y*10**x,list(range(-5,5)),list(random.randint(0,9,siz
6 alpha.extend(list(map(lambda x,y:y*10**x,list(range(-5,5)),list(random.randint(0,9,siz
7 alpha.sort()
8 alpha
```

Out[119]:

```
[0.0,
0,
0.0,
0,
3.0000000000000004e-05,
4e-05,
8e-05,
0.00030000000000000003,
0.00060000000000000001,
0.001,
0.003,
0.007,
0.01,
0.04,
0.07,
0.1,
0.6000000000000001,
2,
4,
6,
10,
30,
200,
700,
800,
3000,
7000,
10000,
50000,
60000]
```

In [120]:

```
1  #removing 0 as value of alpha
2  while True:
3      if 0 in alpha:
4          alpha.pop(0)
5          continue
6      break
7  alpha
8
```

Out[120]:

```
[3.0000000000000004e-05,
4e-05,
8e-05,
0.00030000000000000003,
0.0006000000000000001,
0.001,
0.003,
0.007,
0.01,
0.04,
0.07,
0.1,
0.6000000000000001,
2,
4,
6,
10,
30,
200,
700,
800,
3000,
7000,
10000,
50000,
60000]
```

In [121]:

```

1 from sklearn.naive_bayes import MultinomialNB as MNB
2 from sklearn.model_selection import GridSearchCV
3
4 param={'alpha':alpha}
5 model = MNB(class_prior=[0.5,0.5])
6 gscv=GridSearchCV(model,param,cv=10,verbose=5,scoring='roc_auc',return_train_score=True)
7 gscv.fit(X_train_bow,y_train_bow)

```

Fitting 10 folds for each of 26 candidates, totalling 260 fits

```

[Parallel(n_jobs=-1)]: Using backend LokyBackend with 2 concurrent workers.
[Parallel(n_jobs=-1)]: Done 24 tasks      | elapsed:    2.0s
[Parallel(n_jobs=-1)]: Done 132 tasks    | elapsed:   10.5s
[Parallel(n_jobs=-1)]: Done 260 out of 260 | elapsed:   20.4s finished

```

Out[121]:

```

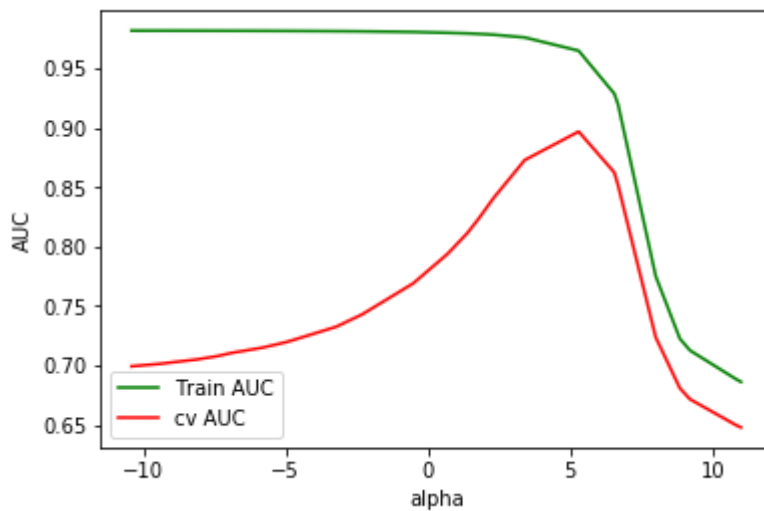
GridSearchCV(cv=10, error_score='raise-deprecating',
             estimator=MultinomialNB(alpha=1.0, class_prior=[0.5, 0.5],
                                     fit_prior=True),
             iid='warn', n_jobs=-1,
             param_grid={'alpha': [3.0000000000000004e-05, 4e-05, 8e-05,
                                   0.00030000000000000003,
                                   0.00060000000000000001, 0.001, 0.003, 0.00
7,
                                   0.01, 0.04, 0.07, 0.1, 0.6000000000000000
1, 2,
                                   4, 6, 10, 30, 200, 700, 800, 3000, 7000,
                                   10000, 50000, 60000]},
             pre_dispatch='2*n_jobs', refit=True, return_train_score=True,
             scoring='roc_auc', verbose=5)

```

In [122]:

```
1 print("Best HyperParameter: ",gscv.best_params_)
2 Best_HyperParameter=gscv.best_params_
3
4 lines=plt.plot(np.log(param['alpha']),gscv.cv_results_['mean_train_score'],np.log(para
5 plt.setp(lines[0],color='g',label='Train AUC')
6 plt.setp(lines[1],color='r',label='cv AUC')
7
8 plt.legend()
9 plt.xlabel("alpha")
10 plt.ylabel("AUC")
11 plt.show()
```

Best HyperParameter: {'alpha': 200}



In [123]:

```

1  from sklearn.metrics import accuracy_score, confusion_matrix, f1_score, precision_score, r
2
3  MODEL = MNB(alpha=Best_HyperParameter['alpha'], class_prior=[0.5, 0.5])
4  MODEL.fit(X_train_bow, y_train_bow)
5  y_pred_bow = MODEL.predict(X_test_bow)
6  y_pred_tr = MODEL.predict(X_train_bow)
7
8  acc=accuracy_score(y_test_bow, y_pred_bow)*100
9  ps=precision_score(y_test_bow, y_pred_bow)*100
10 rc=recall_score(y_test_bow, y_pred_bow)*100
11 f1=f1_score(y_test_bow, y_pred_bow)*100
12
13 print("Accuracy on test set: %0.2f%%"%(acc))
14 print("Precision on test set: %0.2f%%"%(ps))
15 print("recall score on test set: %0.2f%%"%(rc))
16 print("f1 score on test set: %0.2f%%"%(f1))

```

Accuracy on test set: 86.61%

Precision on test set: 97.15%

recall score on test set: 87.05%

f1 score on test set: 91.82%

In [124]:


```

1  cm = pd.DataFrame(confusion_matrix(y_train_bow, y_pred_tr), range(2), range(2))
2  sns.heatmap(cm, annot=True, fmt='g')

```

Out[124]:

<matplotlib.axes._subplots.AxesSubplot at 0x7f967a96e320>




In [125]:

```
1 cm = pd.DataFrame(confusion_matrix(y_test_bow, y_pred_bow), range(2),range(2))
2 sns.heatmap(cm, annot=True,fmt='g')
```

Out[125]:

<matplotlib.axes._subplots.AxesSubplot at 0x7f966ddea90>



In [126]:


```

1  from sklearn.metrics import roc_auc_score
2  train_fpr, train_tpr, thresholds = roc_curve(y_train_bow, MODEL.predict_proba(X_train_
3  test_fpr, test_tpr, thresholds = roc_curve(y_test_bow, MODEL.predict_proba(X_test_bow)
4  print("train auc ={}".format(roc_auc_score(y_train_bow, MODEL.predict_proba(X_train_bo
5  print("test auc=").format(roc_auc_score(y_test_bow, MODEL.predict_proba(X_test_bow)[:
6
7  plt.plot(train_fpr, train_tpr, label="train AUC ")
8  plt.plot(test_fpr, test_tpr, label="train AUC ")
9  plt.legend()
10 plt.xlabel("K: hyperparameter")
11 plt.ylabel("AUC")
12 plt.title("AUC plot")
13 plt.show()

```

train auc =0.9631827341896512

test auc=0.9015188321967882



[5.1.1] Top 10 important features

In [127]:

```
1 neg_sorted = MODEL.feature_log_prob_[0, :].argsort()
2 pos_sorted = MODEL.feature_log_prob_[1, :].argsort()
3
4 print("Important features of Negative class\n")
5 print(np.take(BOW.get_feature_names(), neg_sorted[-10:]))
6 print("\n \n")
7 print("Important features of Positive class \n")
8 print(np.take(BOW.get_feature_names(), pos_sorted[-10:]))
9 print("\n \n")
10
```

Important features of Negative class

```
['would not' 'money' 'disappointed' 'bad' 'one' 'taste' 'product' 'like'
 'would' 'not']
```

Important features of Positive class

```
['product' 'flavor' 'best' 'taste' 'one' 'love' 'like' 'good' 'great'
 'not']
```

[5.2] Applying Naive Bayes on TFIDF, SET 2

In [138]:

```
1 from sklearn.preprocessing import StandardScaler
2 from sklearn.feature_extraction.text import CountVectorizer
3 from sklearn.model_selection import train_test_split
4
5
6 ## splitting with shuffle as false so time consistency is maintained as data is sorted
7
8 X_train_tfidf, X_test_tfidf, y_train_tfidf, y_test_tfidf = train_test_split(preprocess
9
10 tfidf = TfidfVectorizer(ngram_range=(1,2), min_df=10)
11 tfidf.fit_transform(X_train_tfidf)
12 X_train_tfidf = tfidf.transform(X_train_tfidf)
13 X_test_tfidf = tfidf.transform(X_test_tfidf)
14
15 scaler = StandardScaler(with_mean = False)
16 X_train_tfidf = scaler.fit_transform(X_train_tfidf)
17 X_test_tfidf = scaler.transform(X_test_tfidf)
18
19 print(X_train_tfidf.shape)
20 print(X_test_tfidf.shape)
```

(69998, 39325)

(30000, 39325)

In [139]:

```
1 from numpy import random
2 alpha=list()
3 alpha.extend(list(map(lambda x,y:y*10**x,list(range(-5,5)),list(random.randint(0,9,siz
4 alpha.extend(list(map(lambda x,y:y*10**x,list(range(-5,5)),list(random.randint(0,9,siz
5 alpha.extend(list(map(lambda x,y:y*10**x,list(range(-5,5)),list(random.randint(0,9,siz
6 alpha.sort()
7 alpha
```

Out[139]:

```
[0.0,
0.0,
2e-05,
4e-05,
8e-05,
0.0002,
0.00060000000000000001,
0.0008,
0.005,
0.006,
0.01,
0.01,
0.08,
0.30000000000000004,
0.30000000000000004,
1,
2,
5,
10,
50,
50,
300,
500,
800,
4000,
5000,
6000,
20000,
30000,
80000]
```

In [140]:

```
1  #removing 0 as value of alpha
2  while True:
3      if 0 in alpha:
4          alpha.pop(0)
5          continue
6      break
7  alpha
8
```

Out[140]:

```
[2e-05,
4e-05,
8e-05,
0.0002,
0.00060000000000000001,
0.0008,
0.005,
0.006,
0.01,
0.01,
0.08,
0.30000000000000004,
0.30000000000000004,
1,
2,
5,
10,
50,
50,
300,
500,
800,
4000,
5000,
6000,
20000,
30000,
80000]
```

In [141]:

```

1 from sklearn.naive_bayes import MultinomialNB as MNB
2 from sklearn.model_selection import GridSearchCV
3
4 param={'alpha':alpha}
5 model = MNB(class_prior=[0.5,0.5])
6 gscv=GridSearchCV(model,param,cv=10,verbose=5,scoring='roc_auc',return_train_score=True)
7 gscv.fit(X_train_tfidf,y_train_tfidf)

```

Fitting 10 folds for each of 28 candidates, totalling 280 fits

```

[Parallel(n_jobs=-1)]: Using backend LokyBackend with 2 concurrent workers.
[Parallel(n_jobs=-1)]: Done 14 tasks      | elapsed:    1.2s
[Parallel(n_jobs=-1)]: Done 68 tasks      | elapsed:    5.6s
[Parallel(n_jobs=-1)]: Done 158 tasks     | elapsed:   13.0s
[Parallel(n_jobs=-1)]: Done 280 out of 280 | elapsed:   23.1s finished

```

Out[141]:

```

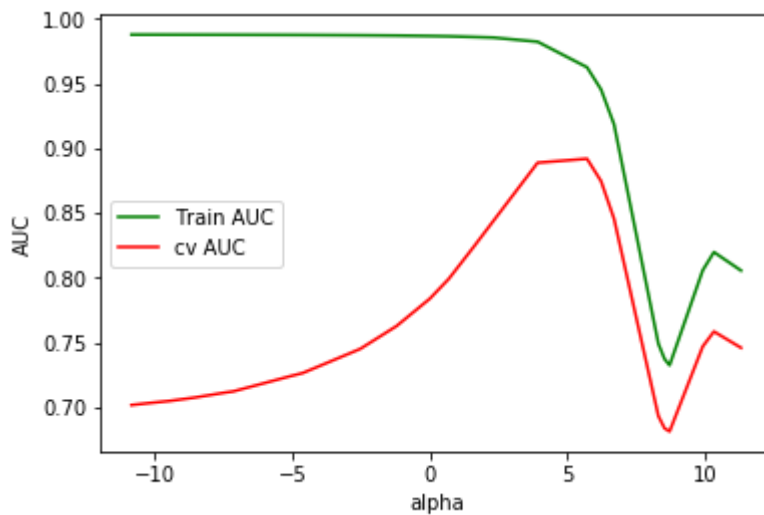
GridSearchCV(cv=10, error_score='raise-deprecating',
             estimator=MultinomialNB(alpha=1.0, class_prior=[0.5, 0.5],
                                     fit_prior=True),
             iid='warn', n_jobs=-1,
             param_grid={'alpha': [2e-05, 4e-05, 8e-05, 0.0002,
                                   0.00060000000000000001, 0.0008, 0.005, 0.0
06,
                                   0.01, 0.01, 0.08, 0.30000000000000004,
                                   0.30000000000000004, 1, 2, 5, 10, 50, 50,
                                   300, 500, 800, 4000, 5000, 6000, 20000,
                                   30000, 80000]},
             pre_dispatch='2*n_jobs', refit=True, return_train_score=True,
             scoring='roc_auc', verbose=5)

```

In [142]:

```
1 print("Best HyperParameter: ",gscv.best_params_)
2 Best_HyperParameter=gscv.best_params_
3
4 lines=plt.plot(np.log(param['alpha']),gscv.cv_results_['mean_train_score'],np.log(para
5 plt.setp(lines[0],color='g',label='Train AUC')
6 plt.setp(lines[1],color='r',label='cv AUC')
7
8 plt.legend()
9 plt.xlabel("alpha")
10 plt.ylabel("AUC")
11 plt.show()
```

Best HyperParameter: {'alpha': 300}



In [143]:

```

1  from sklearn.metrics import accuracy_score, confusion_matrix, f1_score, precision_score, r
2
3  MODEL = MNB(alpha=Best_HyperParameter['alpha'], class_prior=[0.5, 0.5])
4  MODEL.fit(X_train_tfidf, y_train_tfidf)
5  y_pred_tfidf = MODEL.predict(X_test_tfidf)
6  y_pred_tr = MODEL.predict(X_train_tfidf)
7
8  acc=accuracy_score(y_test_tfidf, y_pred_tfidf)*100
9  ps=precision_score(y_test_tfidf, y_pred_tfidf)*100
10 rc=recall_score(y_test_tfidf, y_pred_tfidf)*100
11 f1=f1_score(y_test_tfidf, y_pred_tfidf)*100
12
13 print("Accuracy on test set: %0.2f%%"%(acc))
14 print("Precision on test set: %0.2f%%"%(ps))
15 print("recall score on test set: %0.2f%%"%(rc))
16 print("f1 score on test set: %0.2f%%"%(f1))

```

Accuracy on test set: 86.76%

Precision on test set: 96.77%

recall score on test set: 87.59%

f1 score on test set: 91.95%

In [144]:


```

1  cm = pd.DataFrame(confusion_matrix(y_train_tfidf, y_pred_tr), range(2), range(2))
2  sns.heatmap(cm, annot=True, fmt='g')

```

Out[144]:

<matplotlib.axes._subplots.AxesSubplot at 0x7f9679d575c0>




In [145]:

```
1 cm = pd.DataFrame(confusion_matrix(y_test_tfidf, y_pred_tfidf), range(2),range(2))
2 sns.heatmap(cm, annot=True,fmt='g')
```

Out[145]:

<matplotlib.axes._subplots.AxesSubplot at 0x7f966b064ba8>

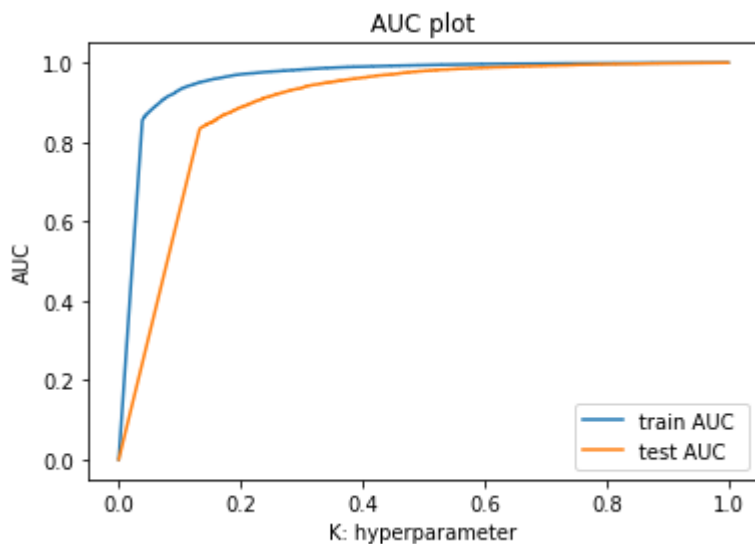


In [146]:

```
1 from sklearn.metrics import roc_auc_score
2 train_fpr, train_tpr, thresholds = roc_curve(y_train_tfidf, MODEL.predict_proba(X_train_tfidf))
3 test_fpr, test_tpr, thresholds = roc_curve(y_test_tfidf, MODEL.predict_proba(X_test_tfidf))
4 print("train auc ={}".format(roc_auc_score(y_train_tfidf, MODEL.predict_proba(X_train_tfidf))))
5 print("test auc ={}".format(roc_auc_score(y_test_tfidf, MODEL.predict_proba(X_test_tfidf))))
6
7 plt.plot(train_fpr, train_tpr, label="train AUC ")
8 plt.plot(test_fpr, test_tpr, label="test AUC ")
9 plt.legend()
10 plt.xlabel("K: hyperparameter")
11 plt.ylabel("AUC")
12 plt.title("AUC plot")
13 plt.show()
```

train auc =0.9619462914727959

test auc=0.893605219887105



[5.2.1] Top 10 important features

In [147]:

```

1 neg_sorted = MODEL.feature_log_prob_[0, :].argsort()
2 pos_sorted = MODEL.feature_log_prob_[1, :].argsort()
3
4 print("Important features of Negative class\n")
5 print(np.take(tfidf.get_feature_names(), neg_sorted[-10:]))
6 print("\n \n")
7 print("Important features of Positive class \n")
8 print(np.take(tfidf.get_feature_names(), pos_sorted[-10:]))
9 print("\n \n")
10

```

Important features of Negative class

```
['money' 'would not' 'disappointed' 'bad' 'one' 'product' 'taste' 'like'
'would' 'not']
```

Important features of Positive class

```
['best' 'product' 'flavor' 'love' 'taste' 'one' 'like' 'good' 'great'
'not']
```

[6] Conclusions

In [154]:

```

1 from prettytable import PrettyTable
2
3 x = PrettyTable()
4
5 x.field_names = ["S.NO", "VECTORIZER", "BEST ALPHA", "TEST AUC", "F1 SCORE"]
6
7 x.add_row(["1", "BOW", 200, 0.9015, 91.82])
8 x.add_row(["2", "TFIDF", 300, 0.8936, 91.95])
9
10 print(x)

```

S.NO	VECTORIZER	BEST ALPHA	TEST AUC	F1 SCORE
1	BOW	200	0.9015	91.82
2	TFIDF	300	0.8936	91.95

1. Naive bayes can be used a benchmark for high dimentional data as being used here.

2. Naive bayes has very low time complexcity

3. BOW vectoriztion is marginaly better than TFIDF vectorization

4. Test AUC and F1 score is even is similar